

8.12 Introduction to Regular Expressions

- A **regular expression** describes a **search pattern** for **matching** characters in other strings
- Can help you extract data from unstructured text
- Can help you ensure that data is in the correct format before processing it

Validating Data

- Often use regular expressions to **validate the data**
 - A U.S. ZIP Code consists of five digits (such as 02215) or five digits followed by a hyphen and four more digits (such as 02215-4775)
 - A string last name contains only letters, spaces, apostrophes and hyphens
 - An e-mail address contains only the allowed characters in the allowed order
 - A U.S. Social Security number contains three digits, a hyphen, two digits, a hyphen and four digits, and adheres to other rules about the specific numbers that can be used in each group of digits
- Rarely need to create your own regular expressions
- Repositories of existing regular expressions that you can copy and use
 - <https://regex101.com>
 - <http://www.regexlib.com>
 - <https://www.regular-expressions.info>

Other Uses of Regular Expressions

- Extract data from text (sometimes known as **scrapping**)
 - e.g., locating all URLs in a web page
 - You might prefer tools like BeautifulSoup, XPath and lxml for this
- Clean data
 - Removing data that's not required, removing duplicate data, handling incomplete data, fixing typos, ensuring consistent data formats, dealing with outliers and more
- Transform data into other formats
 - Reformatting data that was collected as tab-separated or space-separated values into comma-separated values (CSV) for an application that requires data to be in CSV format

©1992–2020 by Pearson Education, Inc. All Rights Reserved. This content is based on Chapter 5 of the book [Intro to Python for Computer Science and Data Science: Learning to Program with AI, Big Data and the Cloud](https://amzn.to/2VvdxnE) (<https://amzn.to/2VvdxnE>).

DISCLAIMER: The authors and publisher of this book have used their best efforts in preparing the book. These efforts include the development, research, and testing of the theories and programs to determine their effectiveness. The authors and publisher make no warranty of any kind, expressed or implied, with regard to these programs or to the documentation contained in these books. The authors and publisher shall not be liable in any event for incidental or consequential damages in connection with, or arising out of, the furnishing, performance, or use of these programs.

8.12.1 re Module and Function fullmatch

- **fullmatch** checks whether the **entire** string in its second argument matches the pattern in its first argument

In [26]:

```
import re
```

Matching Literal Characters

In [27]:

```
pattern = '02215'
```

In [28]:

```
'Match' if re.fullmatch(pattern, '02215') else 'No match'
```

Out[28]:

```
'Match'
```

In [29]:

```
'Match' if re.fullmatch(pattern, '51220') else 'No match'
```

Out[29]:

```
'No match'
```

- First argument is the regular expression pattern to match
 - Any string can be a regular expression
 - **literal characters** match **themselves** in the specified order
- Second argument is the string that should entirely match the pattern.
- If the second argument matches the pattern in the first argument, **fullmatch** returns an object containing the matching text, which evaluates to **True**
- For no match, **fullmatch** returns **None**, which evaluates to **False**

Metacharacters, Character Classes and Quantifiers

- Regular expressions typically contain various special symbols called **metacharacters**:

Regular expression metacharacters

[] { } () \ * + ^ \$? . |

- **\ metacharacter** begins each predefined **character class**
- Each matches a specific set of characters

In [30]:

```
'Valid' if re.fullmatch(r'\d{5}', '02215') else 'Invalid'
```

Out[30]:

```
'Valid'
```

In [31]:

```
'Valid' if re.fullmatch(r'\d{5}', '9876') else 'Invalid'
```

Out[31]:

```
'Invalid'
```

- In `\d{5}`, `\d` is a character class representing a digit (0–9)
- A character class is a **regular expression escape sequence** that matches **one** character
- To match more than one, follow the character class with a **quantifier**
- `{5}` repeats `\d` five times to match five consecutive digits

Other Predefined Character Classes

- To match any metacharacter as its **literal** value, precede it by a backslash (`\`)
 - For example, `\\` matches a backslash (`\`) and `\$` matches a dollar sign (`$`)

Character class	Matches
<code>\d</code>	Any digit (0–9).
<code>\D</code>	Any character that is <i>not</i> a digit.
<code>\s</code>	Any whitespace character (such as spaces, tabs and newlines).
<code>\S</code>	Any character that is <i>not</i> a whitespace character.
<code>\w</code>	Any word character (also called an alphanumeric character)—that is, any uppercase or lowercase letter, any digit or an underscore
<code>\W</code>	Any character that is <i>not</i> a word character.

Custom Character Classes

- Square brackets, `[]`, define a **custom character class** that matches a **single** character
- `[aeiou]` matches a lowercase vowel
- `[A-Z]` matches an uppercase letter
- `[a-z]` matches a lowercase letter
- `[a-zA-Z]` matches any lowercase or uppercase letter

In [32]:

```
'Valid' if re.fullmatch('[A-Z][a-z]*', 'Wally') else 'Invalid'
```

Out[32]:

```
'Valid'
```

In [33]:

```
'Valid' if re.fullmatch('[A-Z][a-z]*', 'eva') else 'Invalid'
```

Out[33]:

```
'Invalid'
```

In [34]:

```
'Match' if re.fullmatch('[^a-z]', 'A') else 'No match'
```

Out[34]:

```
'Match'
```

In [35]:

```
'Match' if re.fullmatch('[^a-z]', 'a') else 'No match'
```

Out[35]:

```
'No match'
```

In [36]:

```
'Match' if re.fullmatch('[*+$$]', '*') else 'No match'
```

Out[36]:

```
'Match'
```

In [37]:

```
'Match' if re.fullmatch('[*+$$]', '!') else 'No match'
```

Out[37]:

```
'No match'
```

* vs. + Quantifier

- + matches **at least one occurrence** of a subexpression
- * and + are **greedy**—they match as many characters as possible

In [38]:

```
'Valid' if re.fullmatch('[A-Z][a-z]+', 'Wally') else 'Invalid'
```

Out[38]:

```
'Valid'
```

In [39]:

```
'Valid' if re.fullmatch('[A-Z][a-z]+', 'E') else 'Invalid'
```

Out[39]:

```
'Invalid'
```

Other Quantifiers

- **? quantifier** matches **zero or one occurrences** of a subexpression

In [40]:

```
'Match' if re.fullmatch('labell?ed', 'labelled') else 'No match'
```

Out[40]:

```
'Match'
```

In [41]:

```
'Match' if re.fullmatch('labell?ed', 'labeled') else 'No match'
```

Out[41]:

```
'Match'
```

In [42]:

```
'Match' if re.fullmatch('labell?ed', 'labelllled') else 'No match'
```

Out[42]:

```
'No match'
```

- `labell?ed` matches `labelled` (the U.K. English spelling) and `labeled` (the U.S. English spelling), but not the misspelled word `labelllled`
- `1?` indicates that there can be **zero or one more** `1` characters before the remaining literal `ed` characters

Other Quantifiers (cont.)

- Can match **at least n occurrences** of a subexpression with the **`{n,}` quantifier**

In [43]:

```
'Match' if re.fullmatch(r'\d{3,}', '123') else 'No match'
```

Out[43]:

```
'Match'
```

In [44]:

```
'Match' if re.fullmatch(r'\d{3,}', '1234567890') else 'No match'
```

Out[44]:

```
'Match'
```

In [45]:

```
'Match' if re.fullmatch(r'\d{3,}', '12') else 'No match'
```

Out[45]:

```
'No match'
```

Other Quantifiers (cont.)

- Can match **between n and m (inclusive) occurrences** of a subexpression with the **$\{n,m\}$ quantifier**

In [46]:

```
'Match' if re.fullmatch(r'\d{3,6}', '123') else 'No match'
```

Out[46]:

```
'Match'
```

In [47]:

```
'Match' if re.fullmatch(r'\d{3,6}', '123456') else 'No match'
```

Out[47]:

```
'Match'
```

In [48]:

```
'Match' if re.fullmatch(r'\d{3,6}', '1234567') else 'No match'
```

Out[48]:

```
'No match'
```

In [49]:

```
'Match' if re.fullmatch(r'\d{3,6}', '12') else 'No match'
```

Out[49]:

```
'No match'
```

©1992–2020 by Pearson Education, Inc. All Rights Reserved. This content is based on Chapter 5 of the book [Intro to Python for Computer Science and Data Science: Learning to Program with AI, Big Data and the Cloud](https://amzn.to/2VvdxnxE) (<https://amzn.to/2VvdxnxE>).

DISCLAIMER: The authors and publisher of this book have used their best efforts in preparing the book. These efforts include the development, research, and testing of the theories and programs to determine their effectiveness. The authors and publisher make no warranty of any kind, expressed or implied, with regard to these programs or to the documentation contained in these books. The authors and publisher shall not be liable in any event for incidental or consequential damages in connection with, or arising out of, the furnishing, performance, or use of these programs.

8.12.2 Replacing Substrings and Splitting Strings

- `sub` function replaces patterns in a string
- `split` function breaks a string into pieces, based on patterns

Function `sub` —Replacing Patterns

- **`sub` function** replaces **all** occurrences of a pattern with the replacement text you specify

In [109]:

```
import re
```

In [110]:

```
re.sub(r'\t', ' ', '1\t2\t3\t4')
```

Out[110]:

```
'1, 2, 3, 4'
```

- Three required arguments:
 - the *pattern to match* (the tab character `'\t'`)
 - the *replacement text* (`' '`) and
 - the *string to be searched* (`'1\t2\t3\t4'`)
- Keyword argument `count` can be used to specify the maximum number of replacements

In [111]:

```
re.sub(r'\t', ' ', '1\t2\t3\t4', count=2)
```

Out[111]:

```
'1, 2, 3\t4'
```

Function `split`

- **`split` function** *tokenizes* a string, using a regular expression to specify the *delimiter*
- Returns a list of strings

In [112]:

```
re.split(r'\s+', '1, 2, 3,4, 5,6,7,8')
```

Out[112]:

```
['1', '2', '3', '4', '5', '6', '7', '8']
```

- Keyword argument `maxsplit` specifies maximum number of splits

In [113]:

```
re.split(r'\s+', '1, 2, 3,4, 5,6,7,8', maxsplit=3)
```

Out[113]:

```
['1', '2', '3', '4, 5,6,7,8']
```

©1992–2020 by Pearson Education, Inc. All Rights Reserved. This content is based on Chapter 5 of the book **[Intro to Python for Computer Science and Data Science: Learning to Program with AI, Big Data and the Cloud](https://amzn.to/2VvdnxE)** (<https://amzn.to/2VvdnxE>).

DISCLAIMER: The authors and publisher of this book have used their best efforts in preparing the book. These efforts include the development, research, and testing of the theories and programs to determine their effectiveness. The authors and publisher make no warranty of any kind, expressed or implied, with regard to these programs or to the documentation contained in these books. The authors and publisher shall not be liable in any event for incidental or consequential damages in connection with, or arising out of, the furnishing, performance, or use of these programs.

8.12.3 Other Search Functions; Accessing Matches

Function `search` — Finding the First Match Anywhere in a String

- `search` looks in a string for the *first* occurrence of a substring that matches a regular expression and returns a **match object** (of type `SRE_Match`) that contains the matching substring
- Match object's `group` method returns that substring

In [114]:

```
import re
```

In [115]:

```
result = re.search('Python', 'Python is fun')
```

In [116]:

```
result.group() if result else 'not found'
```

Out[116]:

```
'Python'
```

- `search` returns `None` if the string does *not* contain the pattern

In [117]:

```
result2 = re.search('fun!', 'Python is fun')
```

In [118]:

```
result2.group() if result2 else 'not found'
```

Out[118]:

```
'not found'
```

- Can search for a match only at the *beginning* of a string with function `match`

Ignoring Case with the Optional `flags` Keyword Argument

- Many `re` module functions receive an optional `flags` keyword argument
- Changes how regular expressions are matched

In [119]:

```
result3 = re.search('Sam', 'SAM WHITE', flags=re.IGNORECASE)
```

In [120]:

```
result3.group() if result3 else 'not found'
```

Out[120]:

'SAM'

Metacharacters That Restrict Matches to the Beginning or End of a String

- **^ metacharacter** at the beginning of a regular expression (and not inside square brackets) is an anchor
- Indicates that the expression matches only the *beginning* of a string

In [121]:

```
result = re.search('^Python', 'Python is fun')
```

In [122]:

```
result.group() if result else 'not found'
```

Out[122]:

'Python'

In [123]:

```
result = re.search('^fun', 'Python is fun')
```

In [124]:

```
result.group() if result else 'not found'
```

Out[124]:

'not found'

- **\$ metacharacter** at the end of a regular expression is an anchor indicating that the expression matches only the *end* of a string

In [125]:

```
result = re.search('Python$', 'Python is fun')
```

In [126]:

```
result.group() if result else 'not found'
```

Out[126]:

'not found'

In [127]:

```
result = re.search('fun$', 'Python is fun')
```

In [128]:

```
result.group() if result else 'not found'
```

Out[128]:

```
'fun'
```

Function `findall` and `finditer` — Finding All Matches in a String

- **`findall`** finds every matching substring in a string
- Returns a list of the matching substrings

In [129]:

```
contact = 'Wally White, Home: 555-555-1234, Work: 555-555-4321'
```

In [130]:

```
re.findall(r'\d{3}-\d{3}-\d{4}', contact)
```

Out[130]:

```
['555-555-1234', '555-555-4321']
```

- **`finditer`** works like `findall`, but returns a lazy *iterable* of match objects

In [131]:

```
for phone in re.finditer(r'\d{3}-\d{3}-\d{4}', contact):  
    print(phone.group())
```

```
555-555-1234
```

```
555-555-4321
```

Capturing Substrings in a Match

- Use **parentheses metacharacters**— (and) —to capture substrings in a match

In [132]:

```
text = 'Charlie Cyan, e-mail: demo1@deitel.com'
```

In [133]:

```
pattern = r'([A-Z][a-z]+ [A-Z][a-z]+), e-mail: (\w+@\w+\.\w{3})'
```

In [134]:

```
result = re.search(pattern, text)
```

- The regular expression specifies two substrings to capture, each denoted by the metacharacters `(` and `)`
- `(` and `)` do *not* affect whether the `pattern` is found in the string `text`
- `match` function returns a match object *only* if the *entire* `pattern` is found in the string `text`
- `match` object's **groups** method returns a tuple of the captured substrings

In [135]:

```
result.groups()
```

Out[135]:

```
('Charlie Cyan', 'demo1@deitel.com')
```

- `match` object's `group` method returns the *entire* match as a single string

In [136]:

```
result.group()
```

Out[136]:

```
'Charlie Cyan, e-mail: demo1@deitel.com'
```

- Access each captured substring by passing an integer to the `group` method

In [137]:

```
result.group(1)
```

Out[137]:

```
'Charlie Cyan'
```

In [138]:

```
result.group(2)
```

Out[138]:

```
'demo1@deitel.com'
```

©1992–2020 by Pearson Education, Inc. All Rights Reserved. This content is based on Chapter 5 of the book [Intro to Python for Computer Science and Data Science: Learning to Program with AI, Big Data and the Cloud](https://amzn.to/2VvdxnE) (<https://amzn.to/2VvdxnE>).

DISCLAIMER: The authors and publisher of this book have used their best efforts in preparing the book. These efforts include the development, research, and testing of the theories and programs to determine their effectiveness. The authors and publisher make no warranty of any kind, expressed or implied, with regard to these programs or to the documentation contained in these books. The authors and publisher shall not be liable in any event for incidental or consequential damages in connection with, or arising out of, the furnishing, performance, or use of these programs.

8.13 Intro to Data Science: Pandas, Regular Expressions and Data Munging

- Data does not always come in forms ready for analysis
- Data could be
 - wrong format
 - incorrect
 - missing
- Data scientists can spend as much as 75% of their time preparing data before they begin their studies
- Called **data munging** or **data wrangling**

8.13 Intro to Data Science: Pandas, Regular Expressions and Data Munging (cont.)

- Two of the most important steps in data munging are *data cleaning* and *transforming data* into optimal formats for database systems and analytics software
- Common data cleaning examples include:
 - deleting observations with missing values,
 - substituting reasonable values for missing values,
 - deleting observations with bad values,
 - substituting reasonable values for bad values,
 - tossing outliers (although sometimes you'll want to keep them),
 - duplicate elimination (although sometimes duplicates are valid),
 - dealing with inconsistent data,
 - and more.

8.13 Intro to Data Science: Pandas, Regular Expressions and Data Munging (cont.)

- Data cleaning is a difficult and messy process where you could easily make bad decisions that would negatively impact your results
- The actions data scientists take can vary per project, be based on the quality and nature of the data and be affected by evolving organization and professional standards
- Common data transformations include:
 - removing unnecessary data and *features* (we'll say more about features in the data science case studies),
 - combining related features,
 - sampling data to obtain a representative subset (we'll see in the data science case studies that *random sampling* is particularly effective for this and we'll say why),
 - standardizing data formats,
 - grouping data,
 - and more.

Cleaning Your Data

- Bad data values and missing values can significantly impact data analysis
- Some data scientists advise against any attempts to insert “reasonable values”
 - Instead, they advocate clearly marking missing data and leaving it up to the data analytics package to handle the issue

Cleaning Your Data (cont.)

- Consider a hospital that records patients’ temperatures (and probably other vital signs) four times per day
- Assume that the data consists of a name and four `float` values, such as

```
[ 'Brown, Sue', 98.6, 98.4, 98.7, 0.0 ]
```

- Patient’s first three recorded temperatures are 99.7, 98.4 and 98.7
- Last temperature was missing and recorded as 0.0, perhaps because the sensor malfunctioned
- Average of the first three values is 98.57, which is close to normal
- If you calculate the average temperature *including* the missing value for which 0.0 was substituted, the average is only 73.93, clearly a questionable result
- Crucial to “get the data right.”
- One way to clean the data is to substitute a *reasonable* value for the missing temperature, such as the average of the patient’s other readings

Data Validation

- `Series` of five-digit ZIP Codes from a dictionary of city-name/five-digit-ZIP-Code key-value pairs
- Intentionally entered an invalid ZIP Code for Miami

In [1]:

```
import pandas as pd
```

In [2]:

```
zips = pd.Series({'Boston': '02215', 'Miami': '3310'})
```

In [3]:

```
zips
```

Out[3]:

```
Boston    02215
Miami      3310
dtype: object
```

Data Validation (cont.)

- The “second column” represents the `Series`’ ZIP Code *values* (from the dictionary’s values)
- The “first column” represents their *indices* (from the dictionary’s keys)
- Can use regular expressions with Pandas to validate data
- The **`str` attribute** of a `Series` provides string-processing and various regular expression methods
- Use the `str` attribute’s **`match` method** to check whether each ZIP Code is valid:

In [4]:

```
zips.str.match(r'\d{5}')
```

Out[4]:

```
Boston      True
Miami       False
dtype: bool
```

- `match` applies the regular expression `\d{5}` to *each* `Series` element
- Returns a new `Series` containing `True` for each valid element

Data Validation (cont.)

- Several ways to deal with invalid data
- One is to catch it at its source and interact with the source to correct the value
 - Not always possible
- In the case of the bad Miami ZIP Code of 3310, we might look for Miami ZIP Codes beginning with 3310
 - There are two— 33101 and 33109
 - We could pick one of those

Data Validation (cont.)

- Sometimes, rather than matching an *entire* value to a pattern, you’ll want to know whether a value contains a *substring* that matches the pattern
- Use method **`contains`** instead of `match`

In [5]:

```
cities = pd.Series(['Boston, MA 02215', 'Miami, FL 33101'])
```

In [6]:

```
cities
```

Out[6]:

```
0    Boston, MA 02215
1    Miami, FL 33101
dtype: object
```

In [7]:

```
cities.str.contains(r' [A-Z]{2} ')
```

Out[7]:

```
0    True
1    True
dtype: bool
```

In [8]:

```
cities.str.match(r' [A-Z]{2} ')
```

Out[8]:

```
0    False
1    False
dtype: bool
```

Reformatting Your Data

- Consider munging data into a different format
- Assume that an application requires U.S. phone numbers in the format ###-###-####
- The phone numbers have been provided to us as 10-digit strings without hyphens

In [9]:

```
contacts = [ ['Mike Green', 'demo1@deitel.com', '5555555555'],
              ['Sue Brown', 'demo2@deitel.com', '5555551234'] ]
```

In [10]:

```
contactsdf = pd.DataFrame(contacts,
                           columns=[ 'Name', 'Email', 'Phone' ])
```

In [11]:

```
contactsdf
```

Out[11]:

	Name	Email	Phone
0	Mike Green	demo1@deitel.com	5555555555
1	Sue Brown	demo2@deitel.com	5555551234

Reformatting Your Data (cont.)

- Munge the data with functional-style programming
- Can *map* the phone numbers to the proper format by calling the `Series` method **map** on the `DataFrame`'s `'Phone'` column
- `map`'s argument is a *function* that receives a value and returns the *mapped* value
- Our function `get_formatted_phone` maps 10 consecutive digits into the format ###-###-####

In [12]:

```
import re
```

In [13]:

```
def get_formatted_phone(value):  
    result = re.fullmatch(r'(\d{3})(\d{3})(\d{4})', value)  
    return '-'.join(result.groups()) if result else value
```

Reformatting Your Data (cont.)

- Regular expression in the block's first statement matches *only* 10 consecutive digits
- Captures substrings containing the first three digits, the next three digits and the last four digits
- `return` statement:
 - If `result` is `None`, returns `value` unmodified
 - Otherwise, calls `result.groups()` to get a tuple containing the captured substrings and pass that tuple to string method `join` to concatenate the elements, separating each from the next with `'-'` to form the mapped phone number

Reformatting Your Data (cont.)

- Series method `map` returns a new Series containing the results of calling its function argument for each value in the column

In [14]:

```
formatted_phone = contactsdf['Phone'].map(get_formatted_phone)
```

In [15]:

```
formatted_phone
```

Out[15]:

```
0    555-555-5555  
1    555-555-1234  
Name: Phone, dtype: object
```

- Once you've confirmed that the data is in the correct format, you can update it in the original DataFrame

In [16]:

```
contactsdf['Phone'] = formatted_phone
```

In [17]:

```
contactsdf
```

Out[17]:

	Name	Email	Phone
0	Mike Green	demo1@deitel.com	555-555-5555
1	Sue Brown	demo2@deitel.com	555-555-1234

©1992–2020 by Pearson Education, Inc. All Rights Reserved. This content is based on Chapter 5 of the book [Intro to Python for Computer Science and Data Science: Learning to Program with AI, Big Data and the Cloud](https://amzn.to/2VvdxnE) (<https://amzn.to/2VvdxnE>).

DISCLAIMER: The authors and publisher of this book have used their best efforts in preparing the book. These efforts include the development, research, and testing of the theories and programs to determine their effectiveness. The authors and publisher make no warranty of any kind, expressed or implied, with regard to these programs or to the documentation contained in these books. The authors and publisher shall not be liable in any event for incidental or consequential damages in connection with, or arising out of, the furnishing, performance, or use of these programs.