

**APACHE KAFKA & FLINK**  
**Real-time fleet monitoring**

Palermo, 18.12.2023

Speaker: Nicolò Paganin





**SMART INDUSTRY – engineering services company and industrial consulting**



# HEADQUARTERS

**SMART INDUSTRY** legal and operational headquarter in Ozzano dell'Emilia (BO)  
operational headquarters in Italy and abroad:

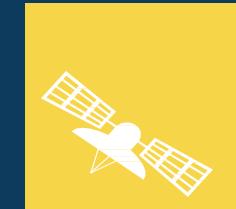
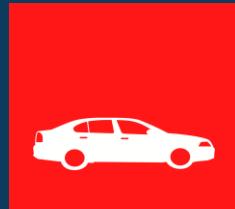


- Trento (TN)
- Torino (TO)
- Amaro (UD)
- Schio (VI)
- Mestre (VE)
- Cernusco sul Naviglio (MI)
- Modena (MO)
- Ozzano dell'Emilia (BO)
- Ancona (AN)
- Roma (RM)
- Napoli (NA)
- Palermo (PA)

- Albania
- Bulgaria
- Romania
- Serbia



# ENGINEERING SECTORS



# STUDENTS OPPORTUNITIES



INTERNSHIPS



DEGREE THESES



WORKSHOPS



Ci occupiamo di **progettazione e sviluppo software**  
In vari ambiti, con varie competenze, in questo caso parleremo di  
**Soluzioni di gestione di grandi moli di dati e stream processing**

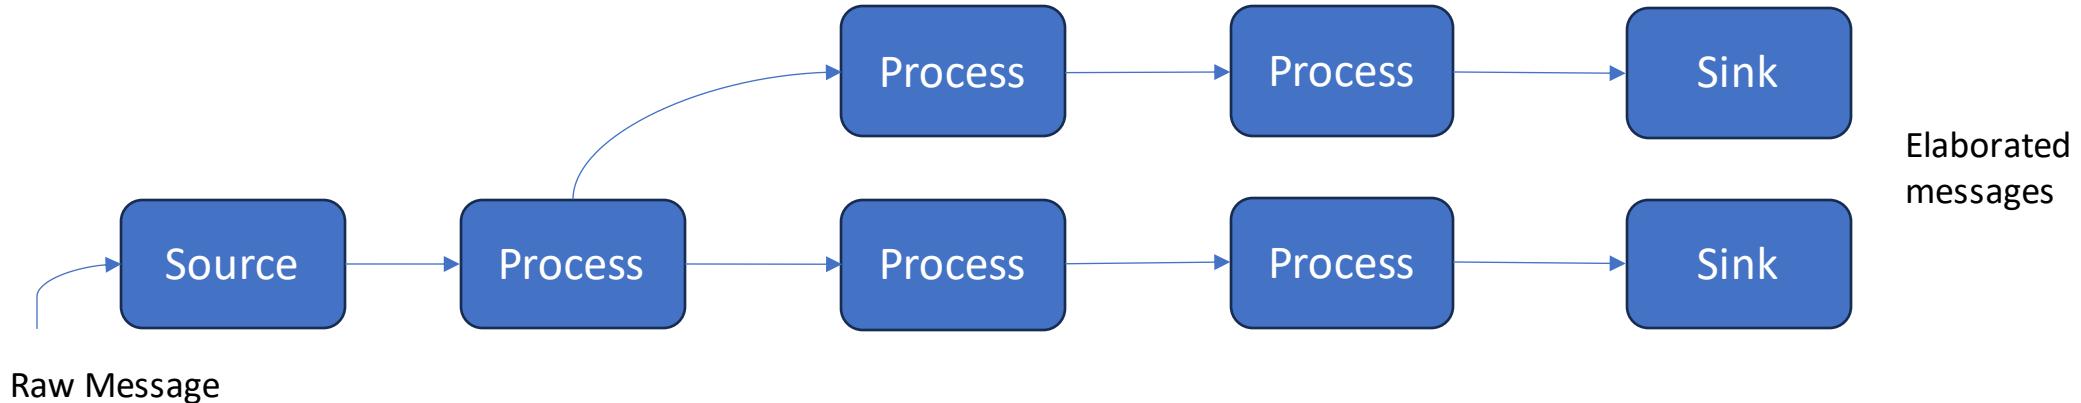
Headquarter: Piazza Maestri del Lavoro 7, Cernusco sul Naviglio (MI)  
Sito web: [www.oncode.it](http://www.oncode.it)

**Speaker: Ing. Nicolò Paganin**  
Managing Partner di Oncode – Direttore BU Enterprise Servizi IT

**Speaker: Andrea Grazioli**  
Technical Team Leader BU Enterprise Servizi IT

# Stream Processing

Stream processing is the elaboration of data in a flow



- Stateless stream processing: elaboration of date does not have a state, so following processing is independent from the past
- Stateful stream processing: elaboration of date has a related state, the state can be changed by current elaboration (es. mean calculation requires a state)



# What we'll talk about

---

Learn about **Apache Kafka (stateless stream processing)** and **Flink (stateful)** and their usage

View a practical application in a car-sharing architecture

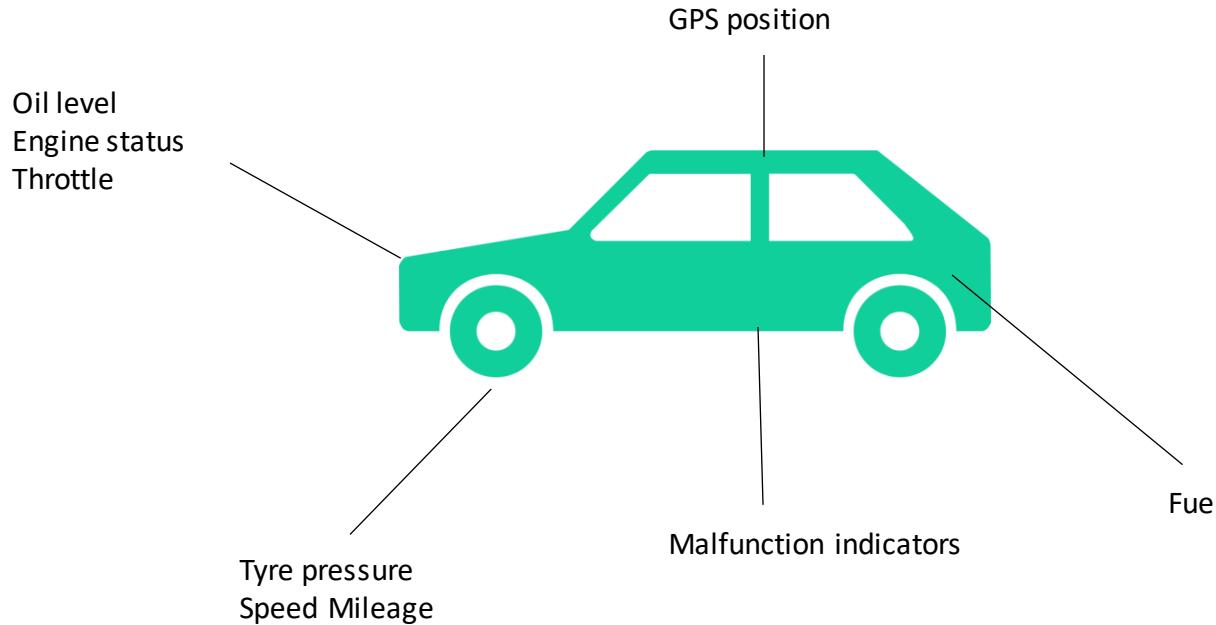
Discover how **car-sharing** services can keep track of their fleets

... and know where you are in  
real time

# Use case

## Data

Even a simple car can be seen as a system that can provide a huge amount of data



# How data are acquired

## Sensors to be installed in the car

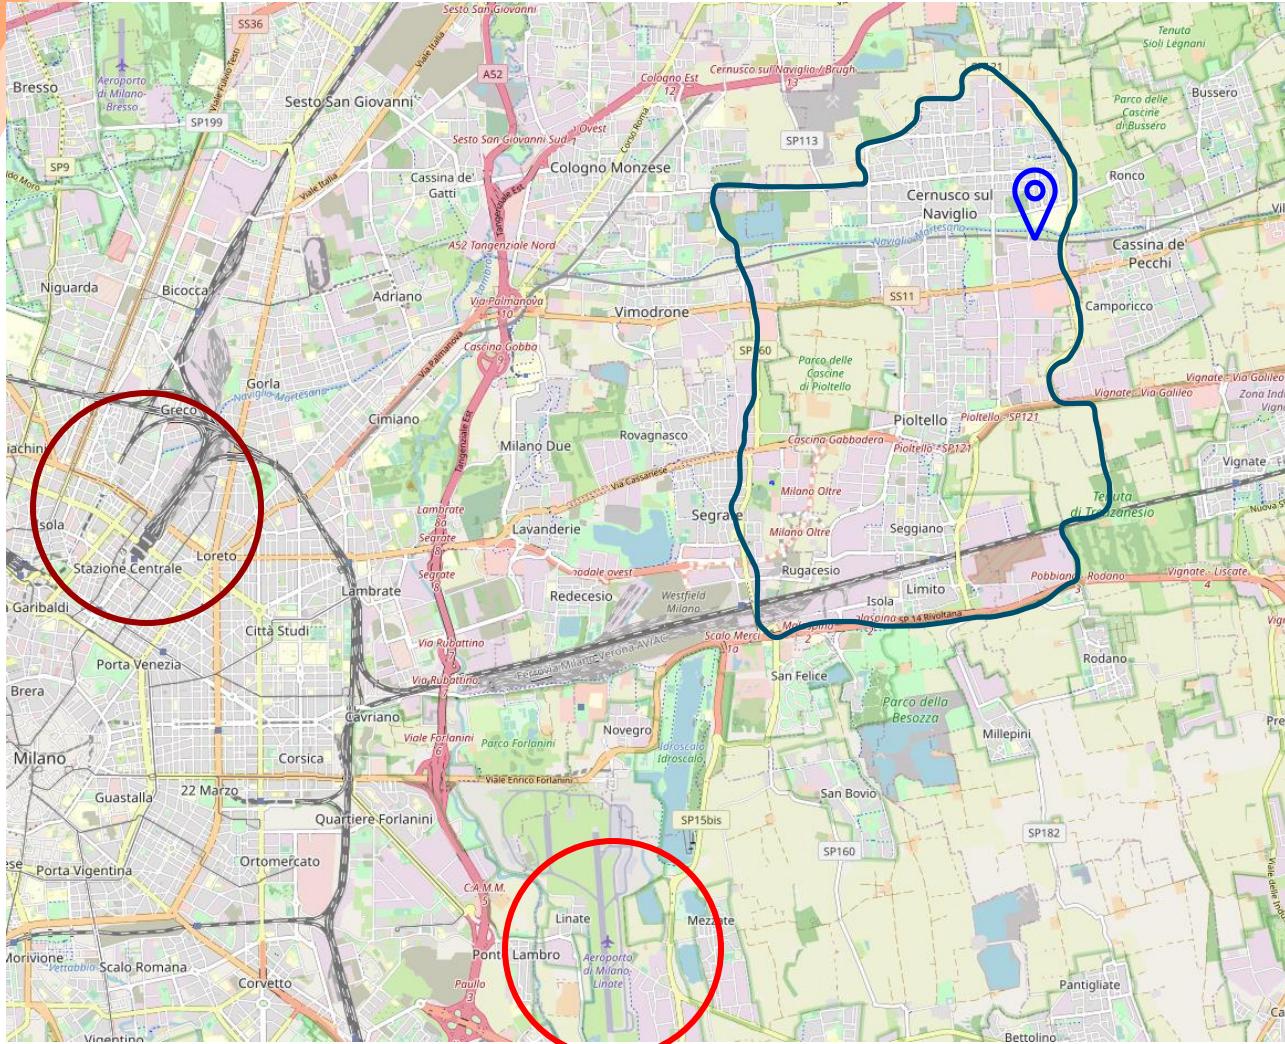


This device collects telemetry data of the car and the GPS location by sending it with a SIM card.

If it is connected via OBD2 to the car  
It may also collect data from the car  
(such as fuel level, car condition,  
central locking, etc. etc.)

# Use case

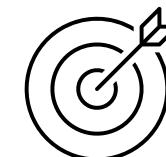
## Geofence



A virtual perimeter for a real-world geographic area such as point of interest (POI), dangerous areas, ZTL.

In this example:

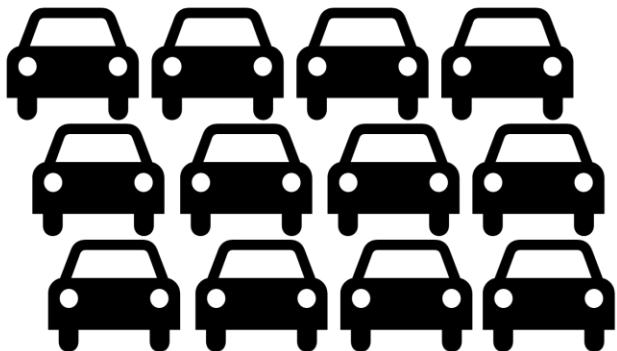
- **Headquarter** 
- **Near office area**
- **Milano Centrale**
- **Linate Airport**



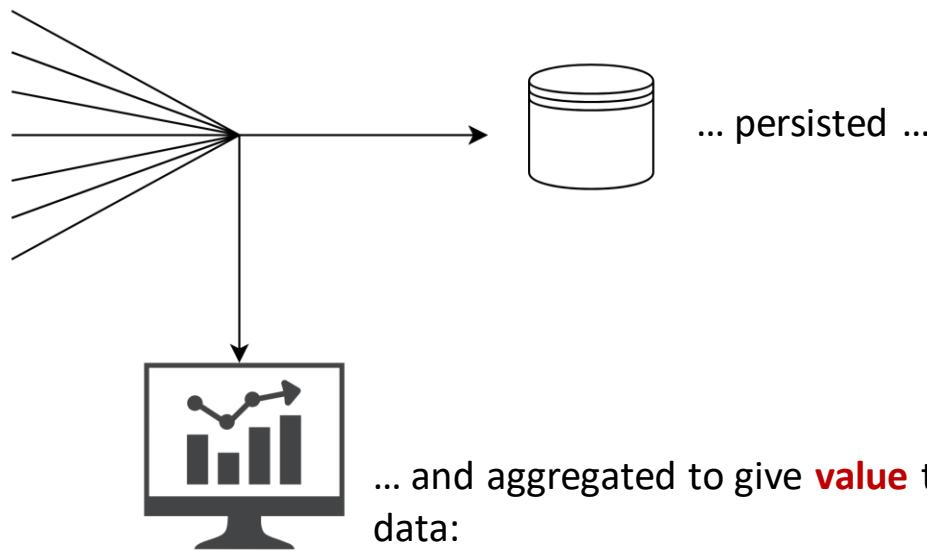
**Detect cars movement through these areas**

# Use case

A simple fleet produces a large amount of such raw data ...



... that raw data must be **collected**, **processed** ...



- ... and aggregated to give **value** to data:
- Keep track of cars
  - Record trips and events

# (Some) Technologies used

---



Log aggregation  
Messaging



**Apache Flink**

Streaming processing  
Complex event detection

# “Collected”: Kafka

---

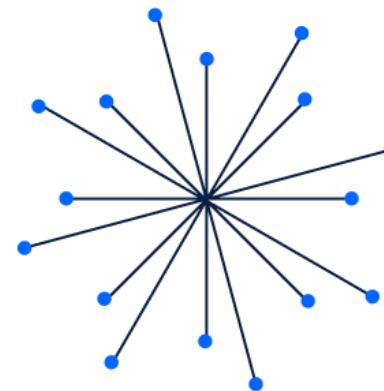


Open-source **distributed event streaming** platform

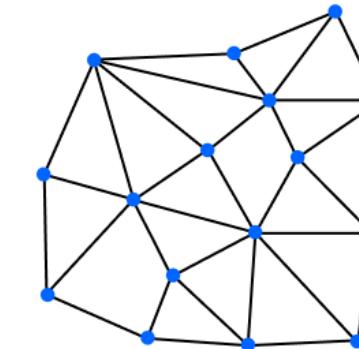
# Apache Kafka

## “Distributed”

Designed to run in clusters, on multiple **scalable** and **reliable** nodes with a set of **servers** and **clients** communicating via TCP



Centralized



Distributed

# Apache Kafka

---

## “Event”

In the data stream, a unit of data represents a specific occurrence (**action, change or state**) in a system.

- Granularity
- Immutable
- Payload
- Event-driven architecture
- Key, value, timestamp, metadata

### Example:

- **Key:** *sensor1*
- **Value:** *pressure 12bar*
- **Timestamp:** *04/12/2023 15:00*

# Apache Kafka

---

## “Streaming”

Apache Kafka manages and processes **data streams** continuously, allowing applications to *read/write* in real-time.

### Data stream:

- Continuous flow of data generated or received over time
- Sequential and ordered
- *Endless*
- Real-time (almost)
- High-volume



# Apache Kafka

---

## Real time streaming

Real time intended as low latency even on large volumes, **NOT** ensuring that a deadline is met!

- **Critical real-time:** microseconds NO
- **Low-latency real-time:** tens/hundreds of microseconds YES
- **Near real-time:** seconds YES

# Apache Kafka

---

## High throughput

Apache Kafka handles large volumes of traffic, ideally scaling up to bandwidth saturation.

- Maximizes sequential writes/reads to **disk**
- Uses an in-memory **cache page**
- Uses a **zero-copy** strategy, reduces redundant copies between cache and buffer

# Apache Kafka

---

## Fault tolerance & Reliability

Apache Kafka provides **reliability**, **consistency** and **redundancy**, thanks to its architecture based on *brokers* and *partitions* (as we will see later).

Kafka is fault tolerant as the failure of a node does not compromise the system thanks to:

- **Redundancy** provided by several replicas (according to the *replication factor* desired)
- A mechanism which elects a new node that acts as leader (*leader election mechanism*), guaranteeing *consistency* for consumers and producers without downtime, errors or interruption

# Apache Kafka

---

## Durability

Messages are written to **filesystems** and cache pages because they offer better performance than in-memory caches or similar structures.

An in-memory cache would take up too much space (28-30GB on a 32GB machine) and would have to be repopulated in the event of a long shutdown (1 GB/minute).

- **Traditional approach:** keep as much as possible in memory, write to disk when it is running out.
- **Kafka approach:** write data to disk, then transfer them to the kernel cache page (pagecache-centric design).

# Apache Kafka

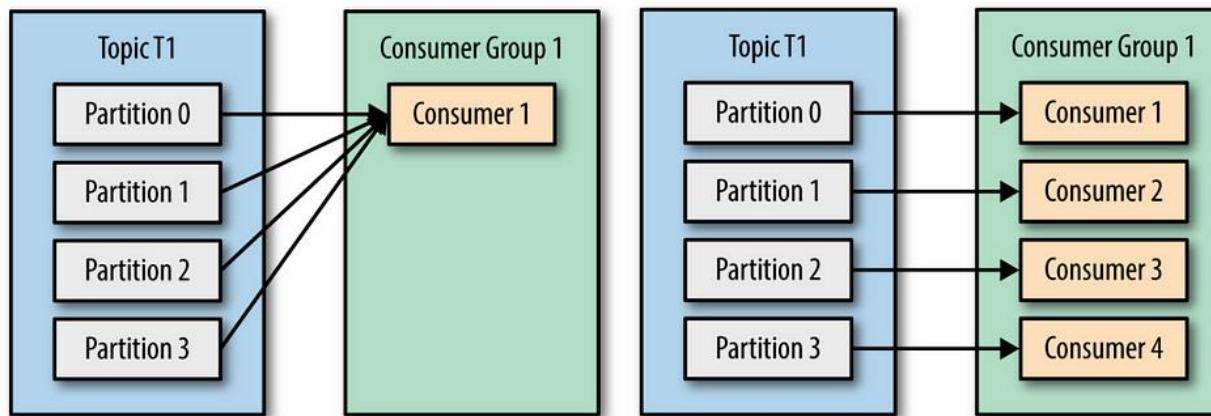
## Scalability

- **Vertically**

Adding more resources to the machine, such as CPU, RAM, disk. More expensive approach.

- **Horizontally**

Adding more brokers to increase processing capacity.  
As we'll see, more is not always better.



# Apache Kafka

---

## General use cases

Apache Kafka can replace traditional message brokers (*like MQTT, AMQP*) to achieve:

- Better throughput
- Partitioning
- Replication
- Fault-tolerance

which are fundamental characteristics for large-scale solutions.

High throughput is often sufficient, but in some use-cases **low end-to-end latency** and **durability** are also useful.

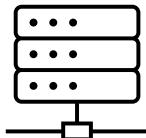
# Apache Kafka

## General use cases



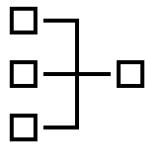
### METRICS

Collect monitoring data and merge them into a centralized system.



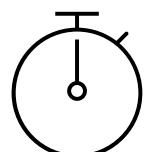
### LOGS

Kafka is often used to aggregate log data, typically from files in a central server.



### STREAM PROCESSING

Processing data in more complex processing pipelines, where raw input data are consumed by Kafka and are subsequently aggregated, enriched or transformed.



### EVENT SOURCING

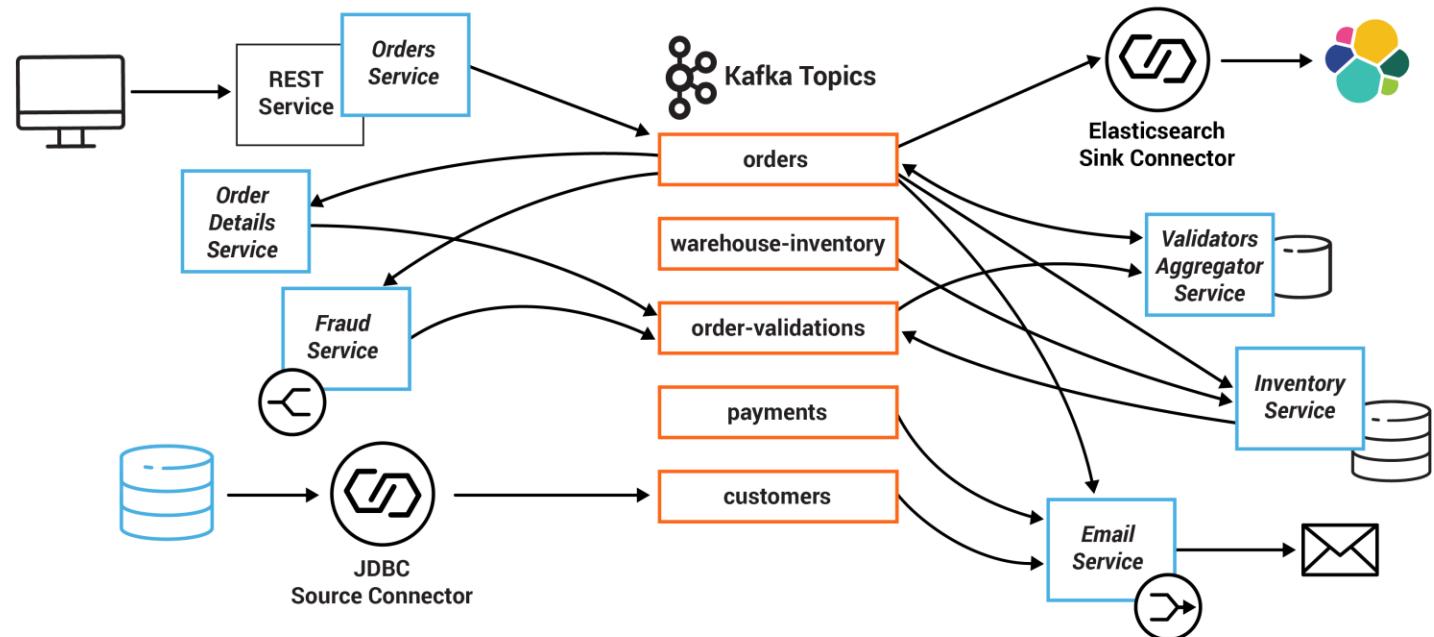
Application design in which state changes are logged as sequences of records in order of time.

BUT ALSO MORE CONCRETELY....

# Apache Kafka

## General use cases

- Processing **transactions** and payments in real-time
- **Monitoring** vehicle fleets
- Reading and **analyzing** data from IoT sensors
- **Responding** in real time to user interactions
- **Connecting** and saving product data between different divisions of a company
- **Connecting** data platforms and microservices

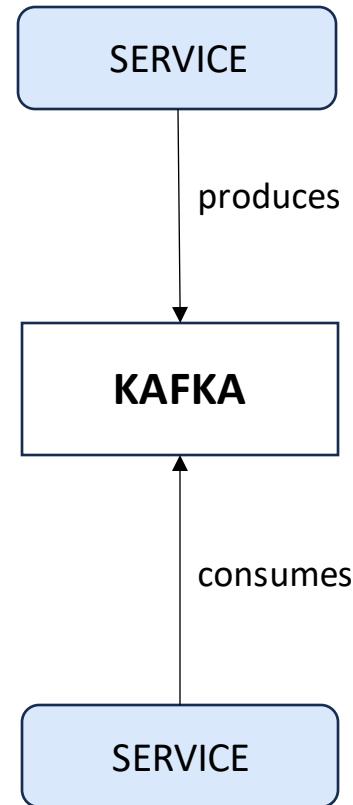


# Apache Kafka Architecture

## Publish/Subscribe messaging

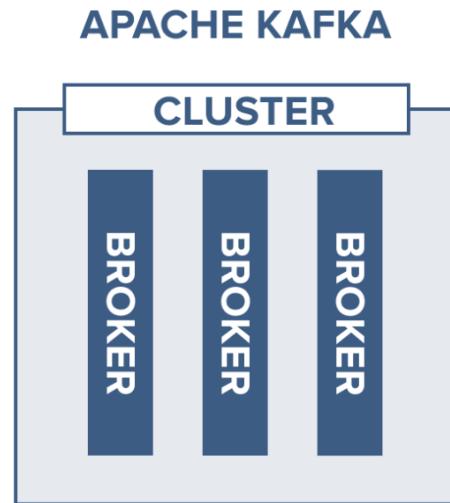
The Pub/Sub messaging architecture:

- Allows to **publish** (write) and **subscribe** (read) on event streams.
- **Decouples** the producer and the consumer of data
- Allows **asynchronous** communication and parallelism



# Apache Kafka Architecture

## Broker & Cluster



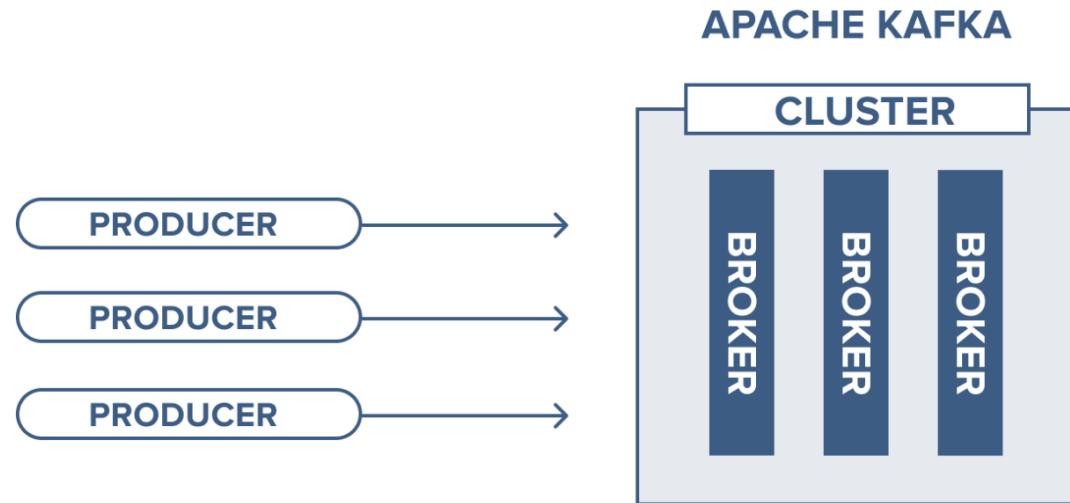
The **cluster**:

- Acts as a **server**
- Manages persistence, replication, durability
- Is scalable
- Maintains Kafka metadata
- Is identified by ID

A single **broker** acts as controller

# Apache Kafka Architecture

## Producer



A **producer** is any application or service that **sends** messages.

It sends directly to the broker who is the leader for that partition.

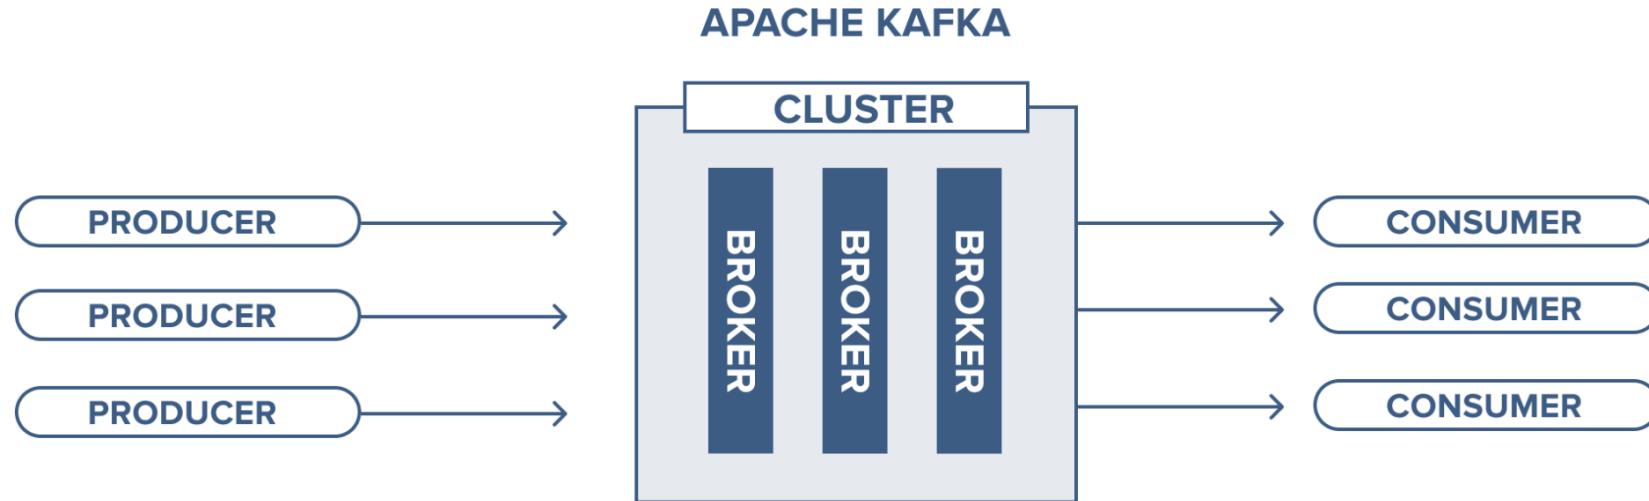
To help it, all nodes respond to a request with metadata about which servers are running and where the leader is located.

Finally, the client checks in which partition it is public.

A producer can also send data as **batch**.

# Apache Kafka Architecture

## Consumer



A **consumer** is any application or service that **receives** messages.

Kafka uses **pull** paradigm, to let the consumer consume at the highest possible rate.

# Apache Kafka Architecture

## Consumer

A key feature of message brokers is **keeping track** of what has been consumed.  
And in Kafka?

- If the broker marks a message as “**consumed**” as soon as it leaves its network?

If the consumer fails to receive it, **the message is lost.**

- If we add ack from the consumer?

Not working if consumer receives but **fails** before ack.

The broker **must keep track** of multiple states for each message

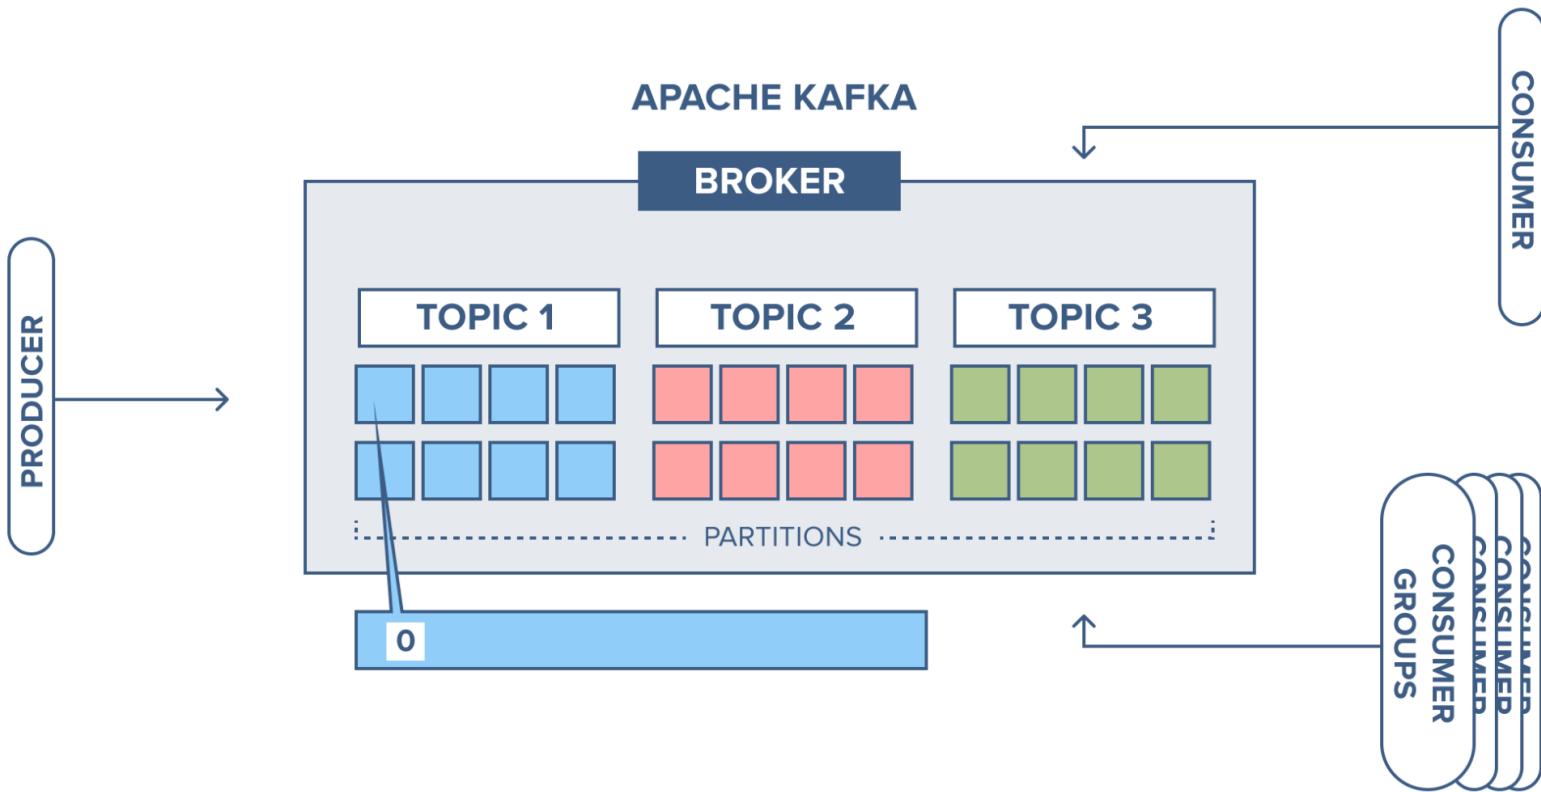
We will see **offset strategy**

# Apache Kafka Architecture

## Topic

A message “category”

Producer → who sends  
Consumer → who reads  
Topic → the subject

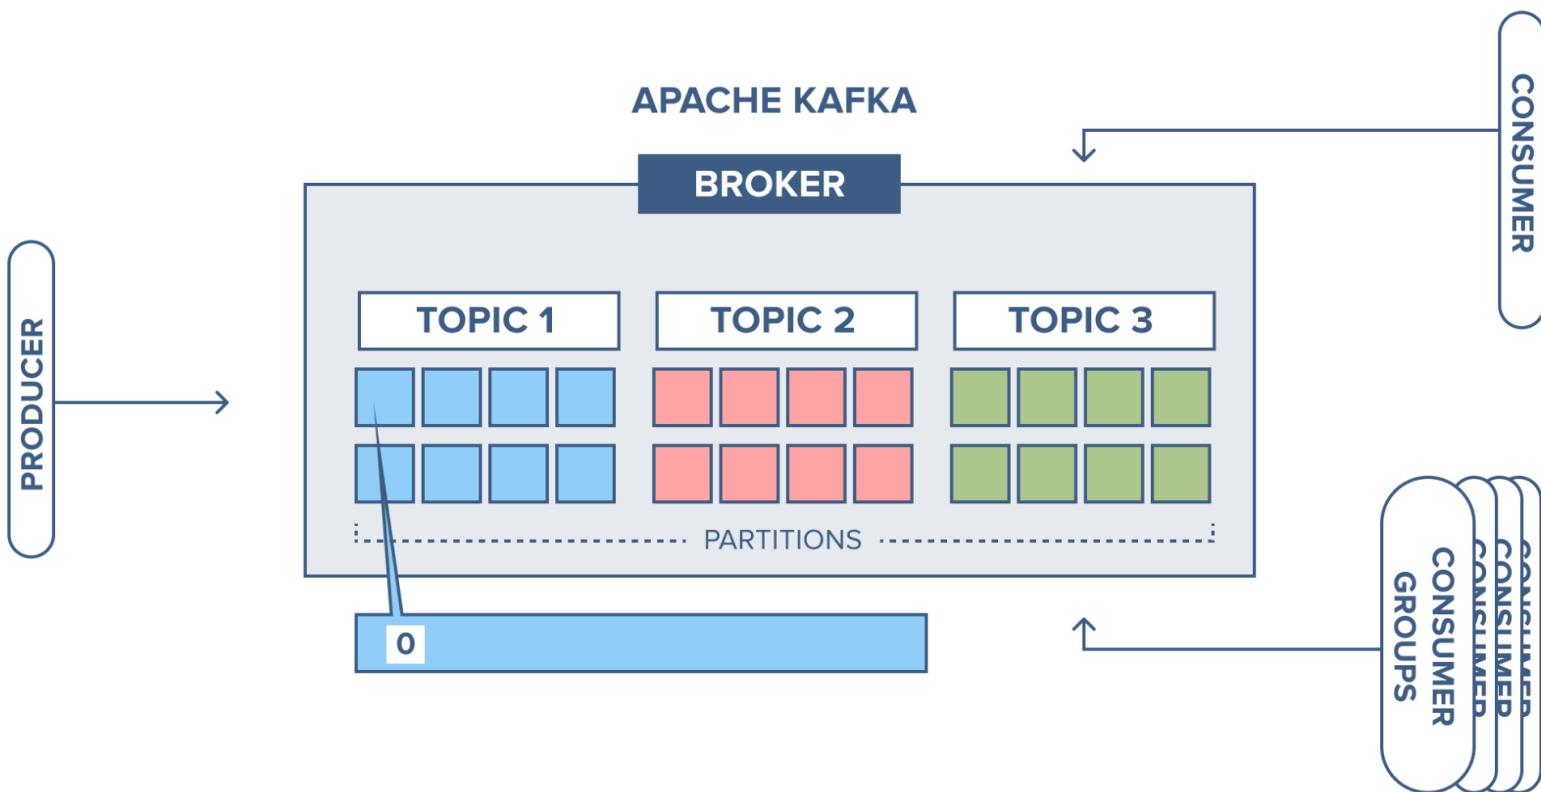


# Apache Kafka Architecture

## Partitions

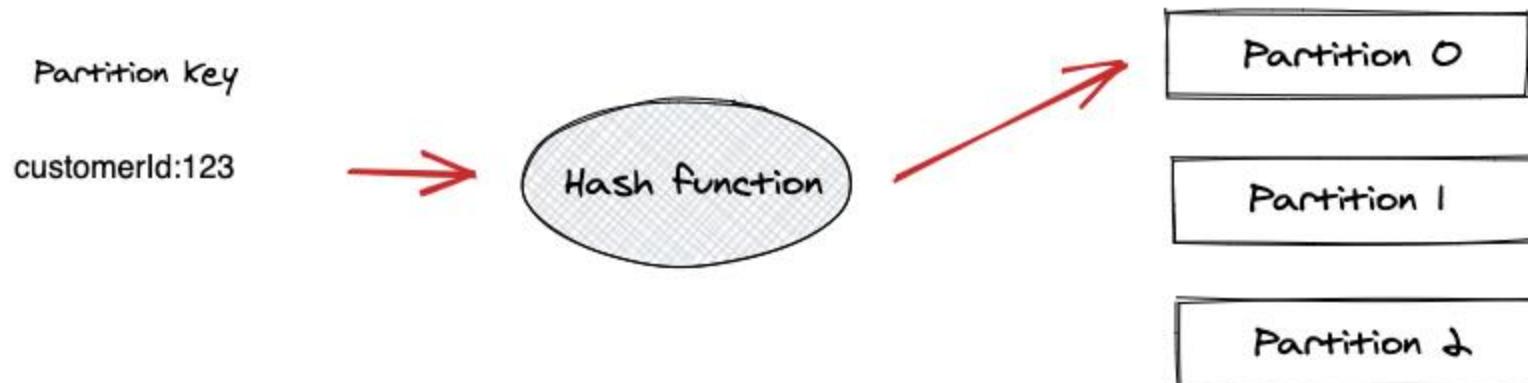
A **partition** is an ordered set of messages with a unique sequence number called **offset**.

Partitions allow greater scalability and fault tolerance



# Apache Kafka Architecture

## Partition key



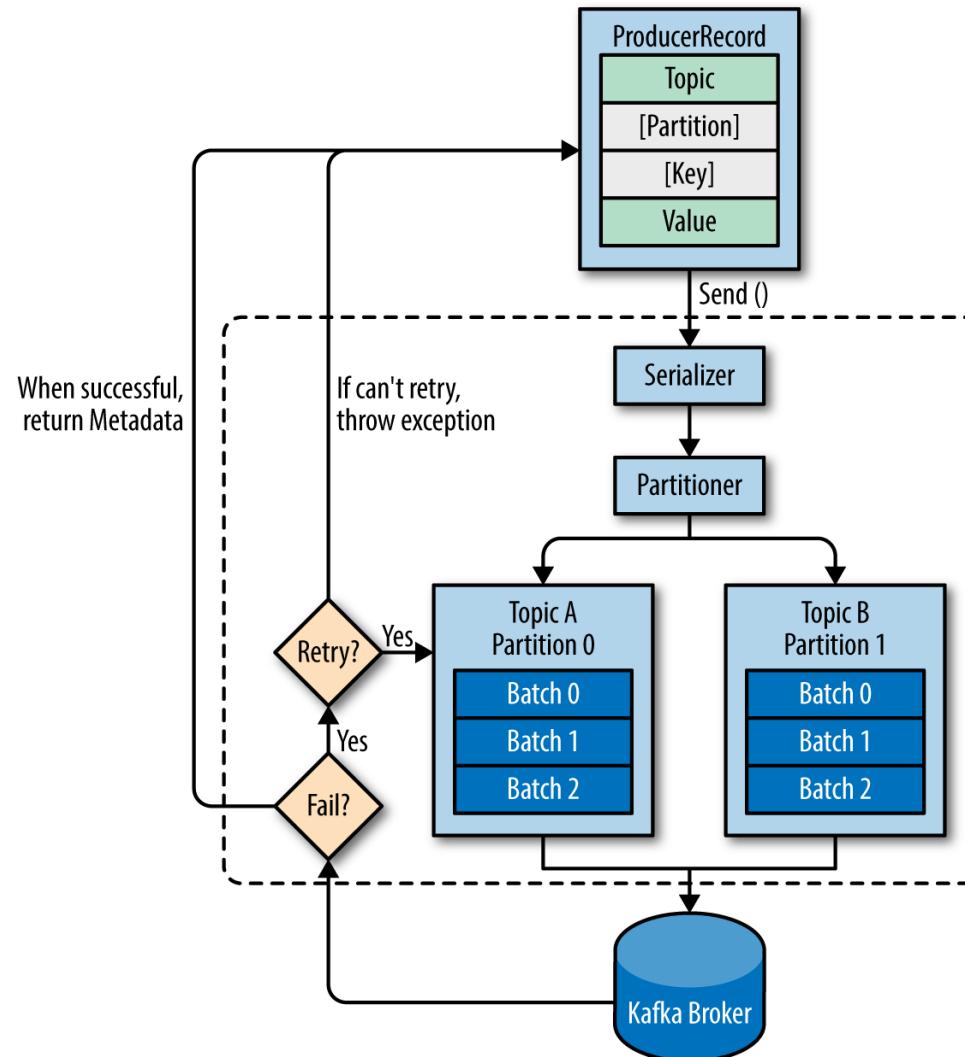
You can choose a partition key to distribute messages.

**NB: choose well!**

- Default Kafka: If not specified, Kafka uses round-robin
- Custom partitioner: implement your own rules

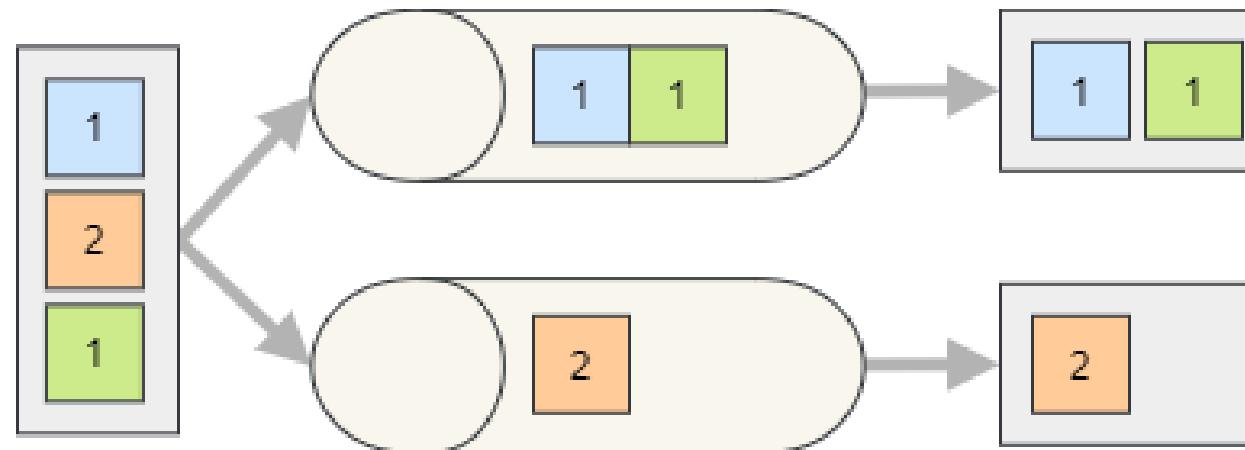
# Apache Kafka Architecture

## Partition key



# Apache Kafka Architecture

## Partition key



Kafka ensures that messages with the same key are written to the same partition

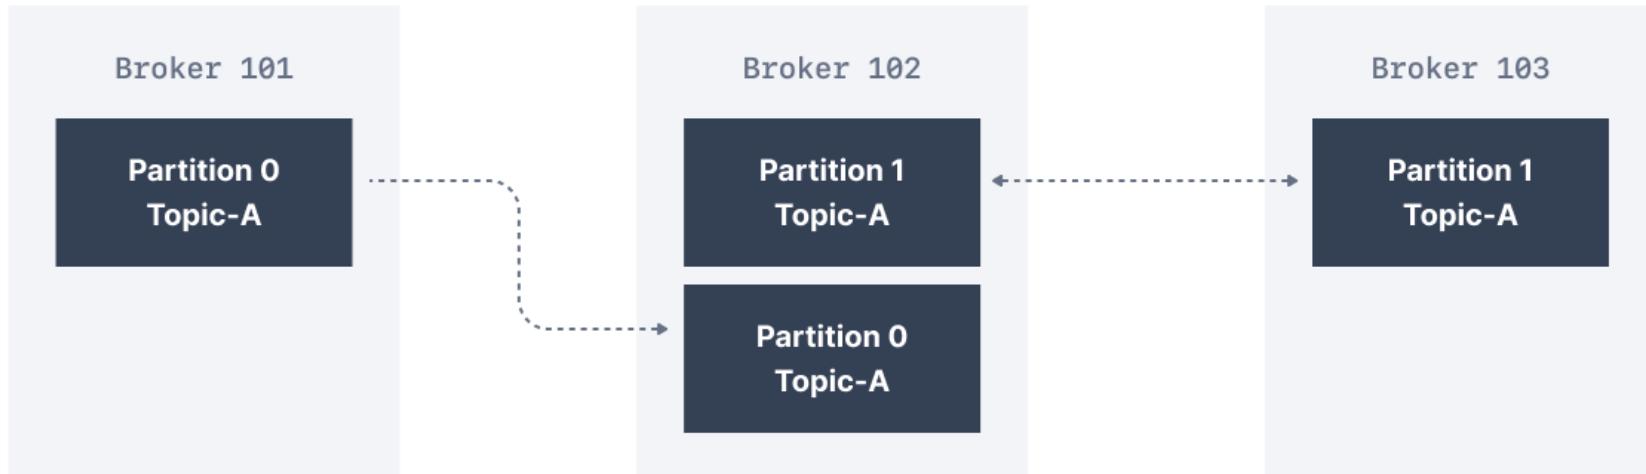
# Apache Kafka

## Replication factor

Factor that regulates how much data replication there will be between brokers.

Prevents data loss and preserves HA (*high availability*) in the event of broker failure.

*Example with Topic-A and replication factor = 2*



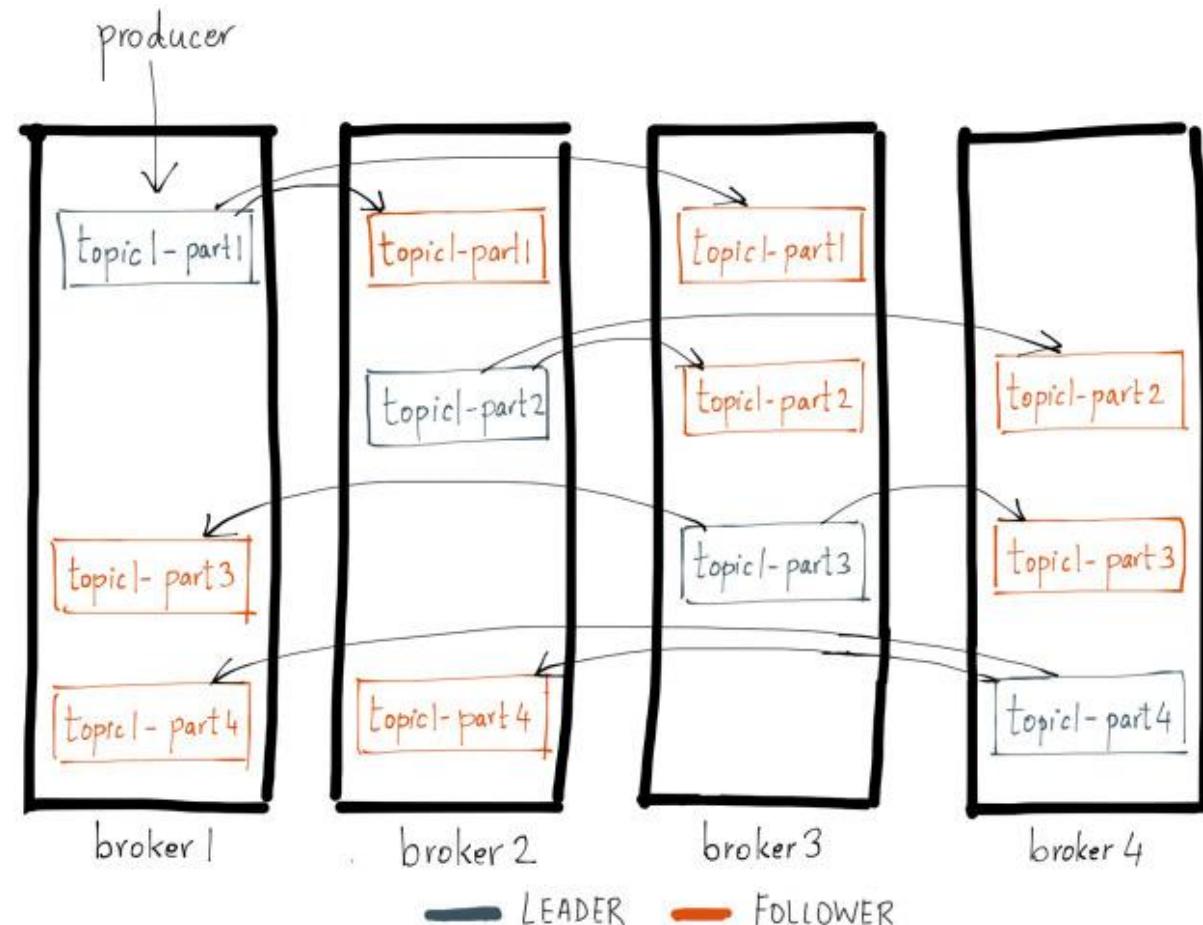
# Apache Kafka

## Replica leader

**Leader** is one of the replicas of a topic partition that is responsible for managing the writing and reading of messages.

It is the **contact point** for producers and consumers.

**Follower** replicas copy the data.



# Apache Kafka

---

## Leader election

The **leader** is fundamental!

Elected by:

- Network latency
- Broker workload
- Specific configurations

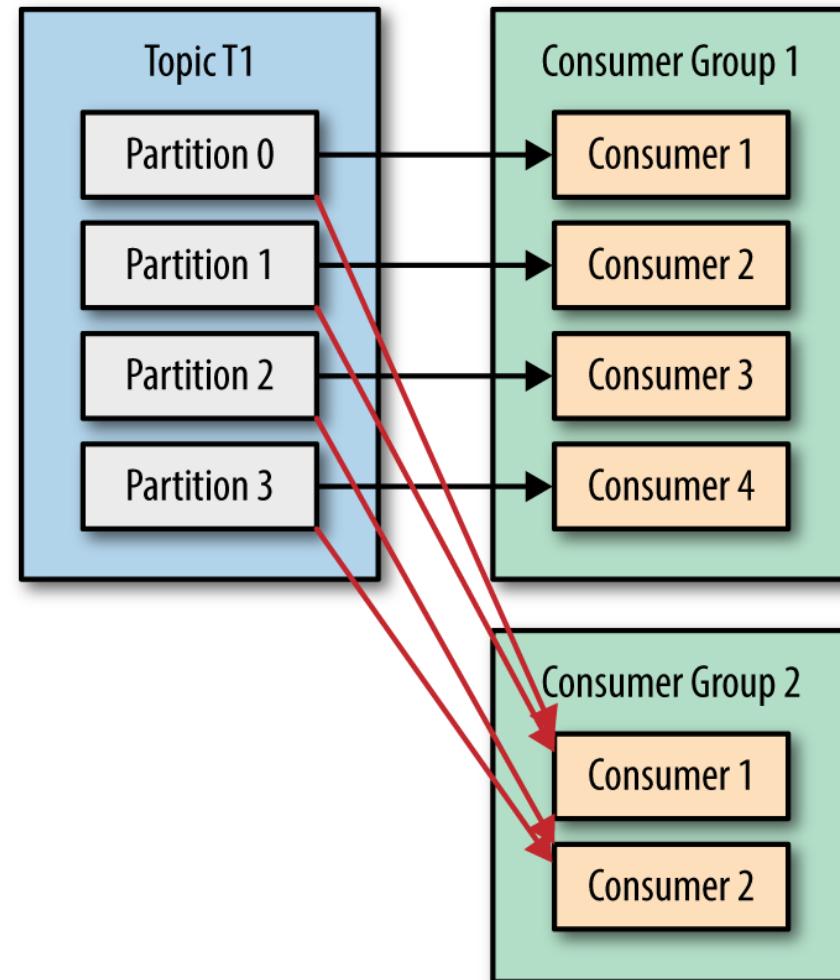
But also:

- Quorum consensus: determined by *replication factor*
- ZooKeeper or topic metadata
- Epoch number
- ISR (In-Sync Replica)
- Performance

# Apache Kafka Architecture

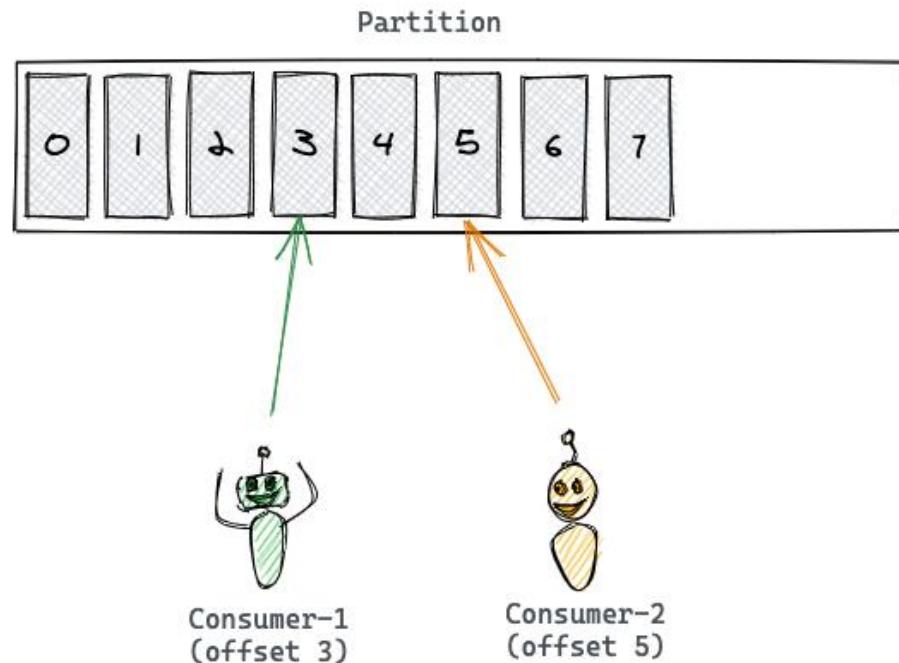
## Consumer Group

- Logical group of consumers united to read messages from one or more topics.
- Unique name
- They allow you to distribute the load among various consumers
- They guarantee that messages are read only **once (per topic, per partition)**



# Apache Kafka Architecture

## Offset



Kafka does not send messages, but consumers read them from the topic.

**Remembering the offset is therefore the consumer's task.**

# Apache Kafka Architecture

---

## Offset

Each consumer has his own “**view**” of the partition and the topic and can **independently** manage the consumption strategy:

- Starting again from the last offset you saved
- Always reprocessing all messages from the beginning
- Starting from a specific offset

# Apache Kafka Architecture

---

## Acks - Producer

The producer sends a message to a Kafka broker and receives an asynchronous ack when the broker has confirmed writing the message to the log.

Using the acks you choose the level of reliability between **at most once, at least once and exactly once**:

1. **acks=0**: the producer does not wait for confirmation, proceeds to the next message
2. **acks=1**: Default, the producer considers a message delivered when leader sends an acknowledgment. **It does not mean that the consumer has received!**
3. **acks=all or acks=-1**: the producer interprets a message as delivered only when all replicating brokers have received the message and written to the logs. More latency!

# Apache Kafka Architecture

---

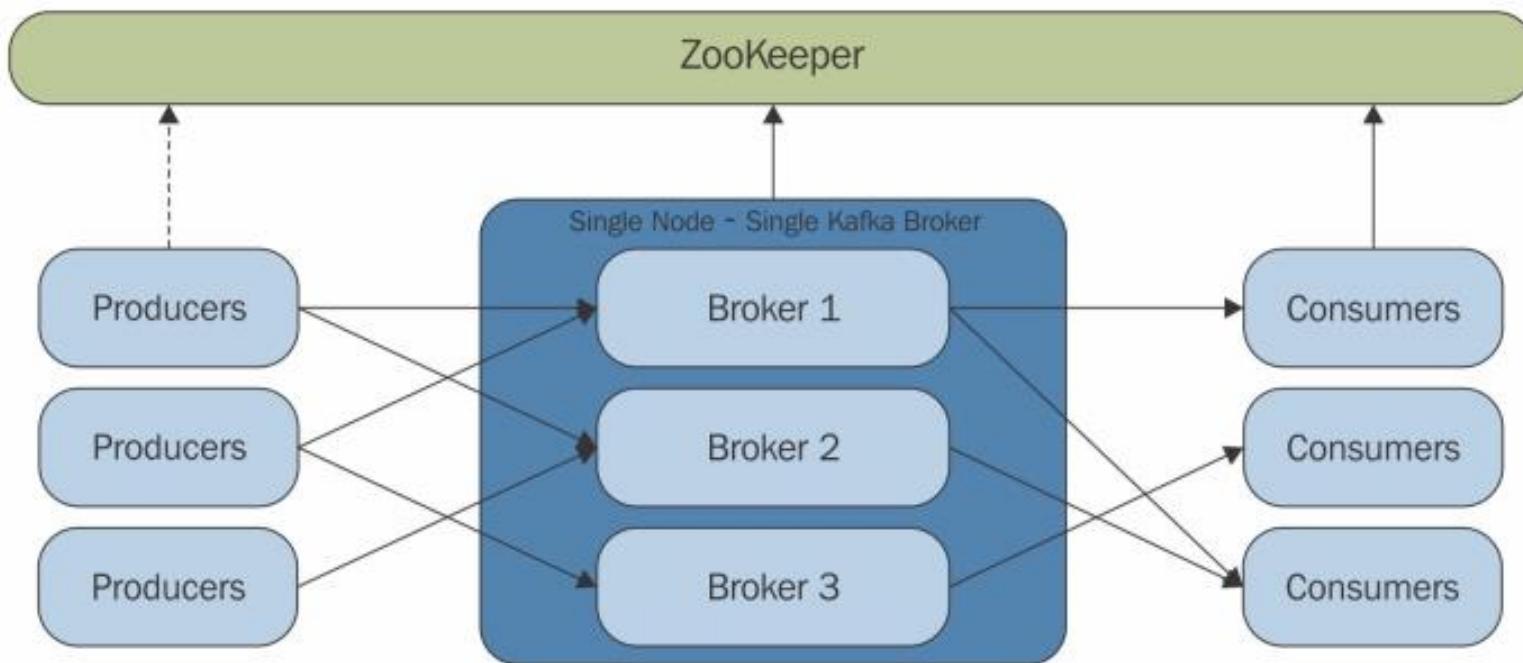
## Commit - Consumer

The consumer can acknowledge receipt of the message and confirm its correct processing.

1. **Auto commit:** managed by Kafka automatically when the message has been delivered to the consumer. Configurable time and frequency.
2. **Explicit commit:** sent manually by the consumer, more control. For example, I await the actual processing of the message.

# Apache Kafka Architecture

## ZooKeeper



Centralized service used to maintain definitions and configurations in a synchronized and flexible way within a distributed system (such as Kafka between brokers).



# Apache Kafka Architecture

---

## ZooKeeper

**ZooKeeper is useful BUT:**

- It adds an extra layer of management
- It makes Kafka more complicated, having to manage two "services" instead of one  
*(Instead, other distributed systems have similar mechanisms but internally managed)*
- It has few hardware needs, but they are still present

# Apache Kafka

---

## Exercises

The repo: <https://tinyurl.com/fleetMonitor>

Who has Docker Desktop installed

Can run the command

docker compose up

to start the containers

In order to compile the Java project and execute use  
this command

`mvn compile exec:java`



# Apache Kafka

---

## Exercise 1

- Create a topic “speeds”
- Write a producer that sends simple random numeric speed values between 0 and 150, for ten random cars every 250 ms.
- JSON Example:

```
{  
  "car": 1,  
  "speed": 95  
}
```

# Apache Kafka

## Exercise 1

```
public static void main(String[] args) {
    try (var producer = createKafkaProducer()) {
        var objectMapper = new ObjectMapper();

        while (true) {
            var random = new Random();
            var vin = random.nextInt(11);
            var speed = random.nextInt(151);

            CarSpeed carSpeed = new CarSpeed(vin, speed);
            String carSpeedJson;

            try {
                carSpeedJson = objectMapper.writeValueAsString(carSpeed);
                Thread.sleep(250);
            } catch (Exception e) {
                throw new RuntimeException(e);
            }
            logger.info("Sending: {}", carSpeedJson);

            ProducerRecord<String, String> record = new ProducerRecord<>(TOPIC, carSpeedJson);
            producer.send(record);
            producer.flush();
        }
    }
}
```

# Apache Kafka

---

## Exercise 1b

- Write a consumer that reads message and print them with metadata (*key, value, offset, partition ecc.*)

# Apache Kafka

## Exercise 1b

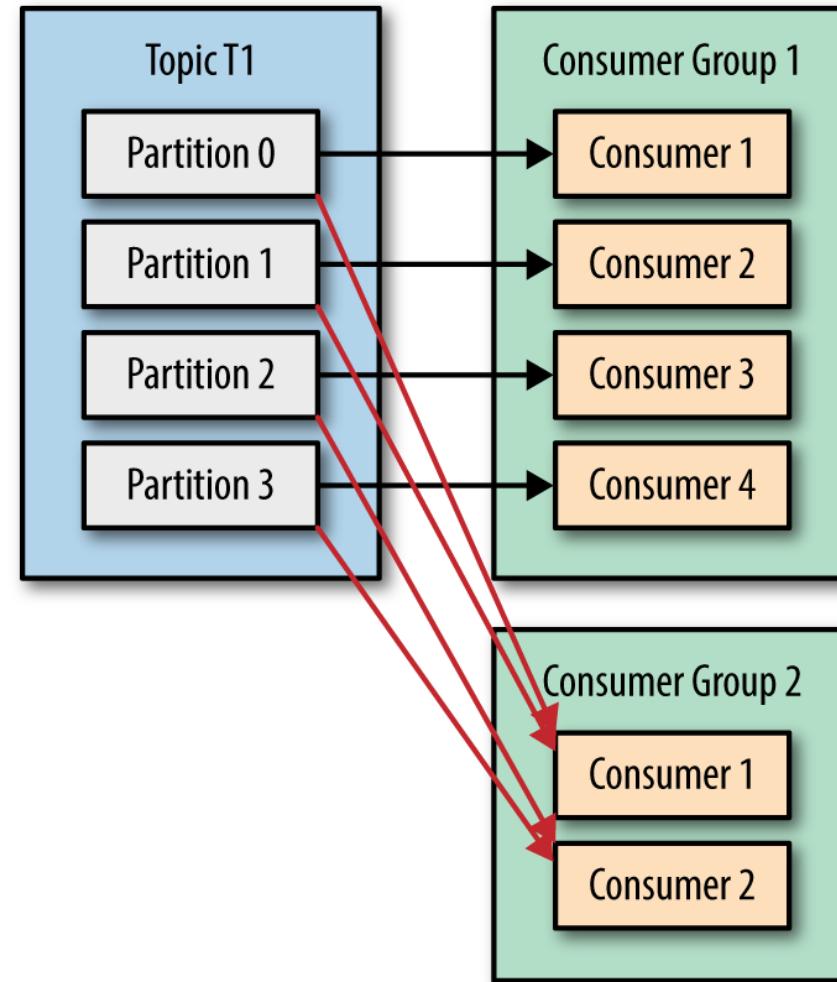
```
public static void main(String[] args) {
    try (var consumer = createKafkaConsumer(consumerGroup)) {
        consumer.subscribe(List.of(topic));

        while (true) {
            ConsumerRecords<String, String> records = consumer.poll(Duration.ofSeconds(1));
            for (ConsumerRecord<String, String> record : records) {
                logger.info("offset = {}, key = {}, value = {}", record.offset(), record.key(), record.value());
            }
        }
    }
}
```

# Apache Kafka Architecture

## Consumer Group

- Logical group of consumers united to read messages from one or more topics.
- Unique name
- They allow you to distribute the load among various consumers
- They guarantee that messages are read only **once (per topic, per partition)**

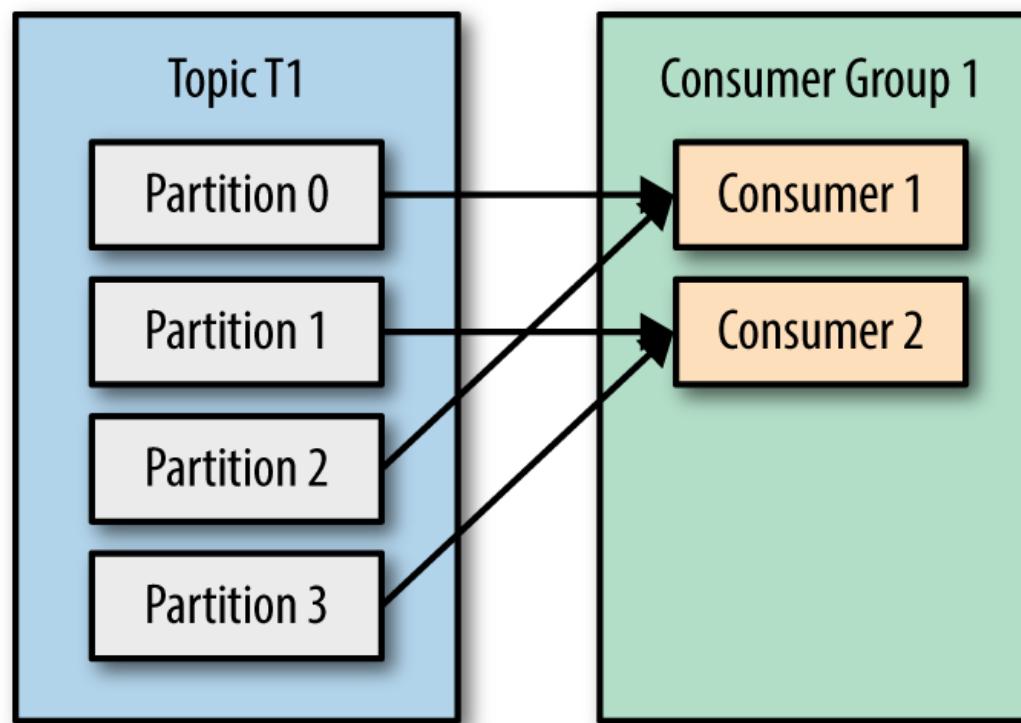


# Apache Kafka

## Parallelism

**P = number of partitions, C = number of consumers**

$P > C$

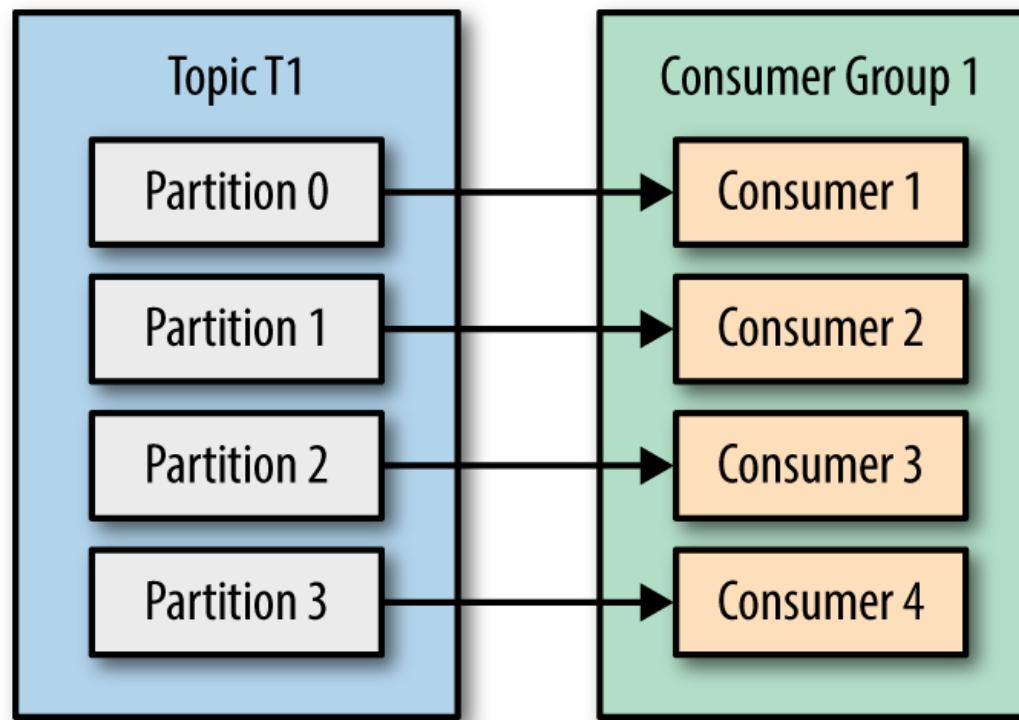


# Apache Kafka

## Parallelism

$P = \text{number of partitions}$ ,  $C = \text{number of consumers}$

$$P = C$$

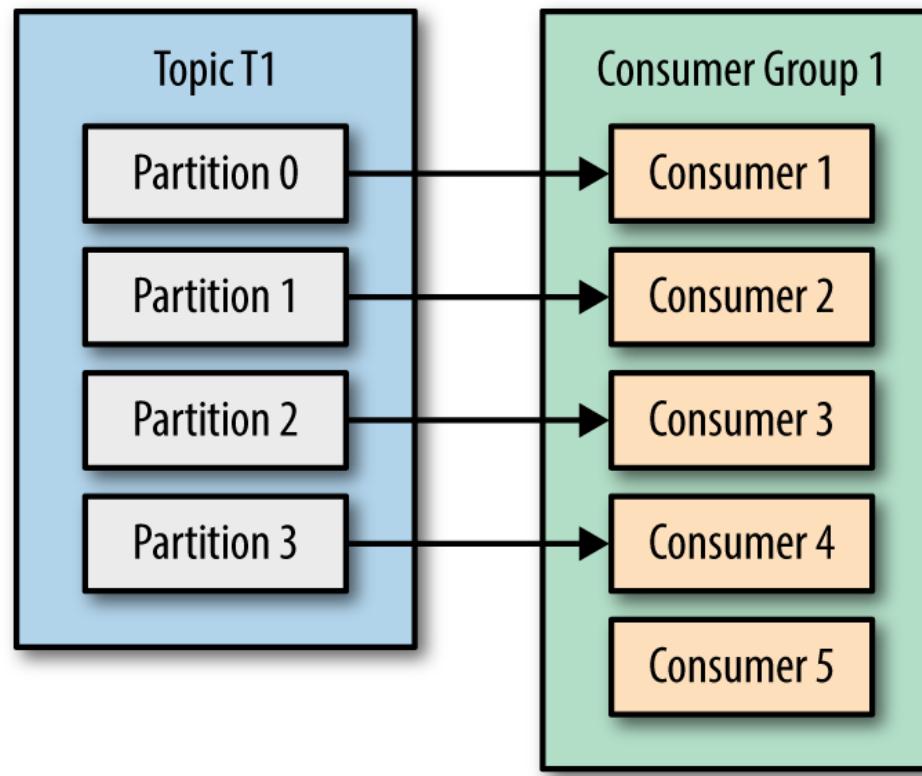


# Apache Kafka

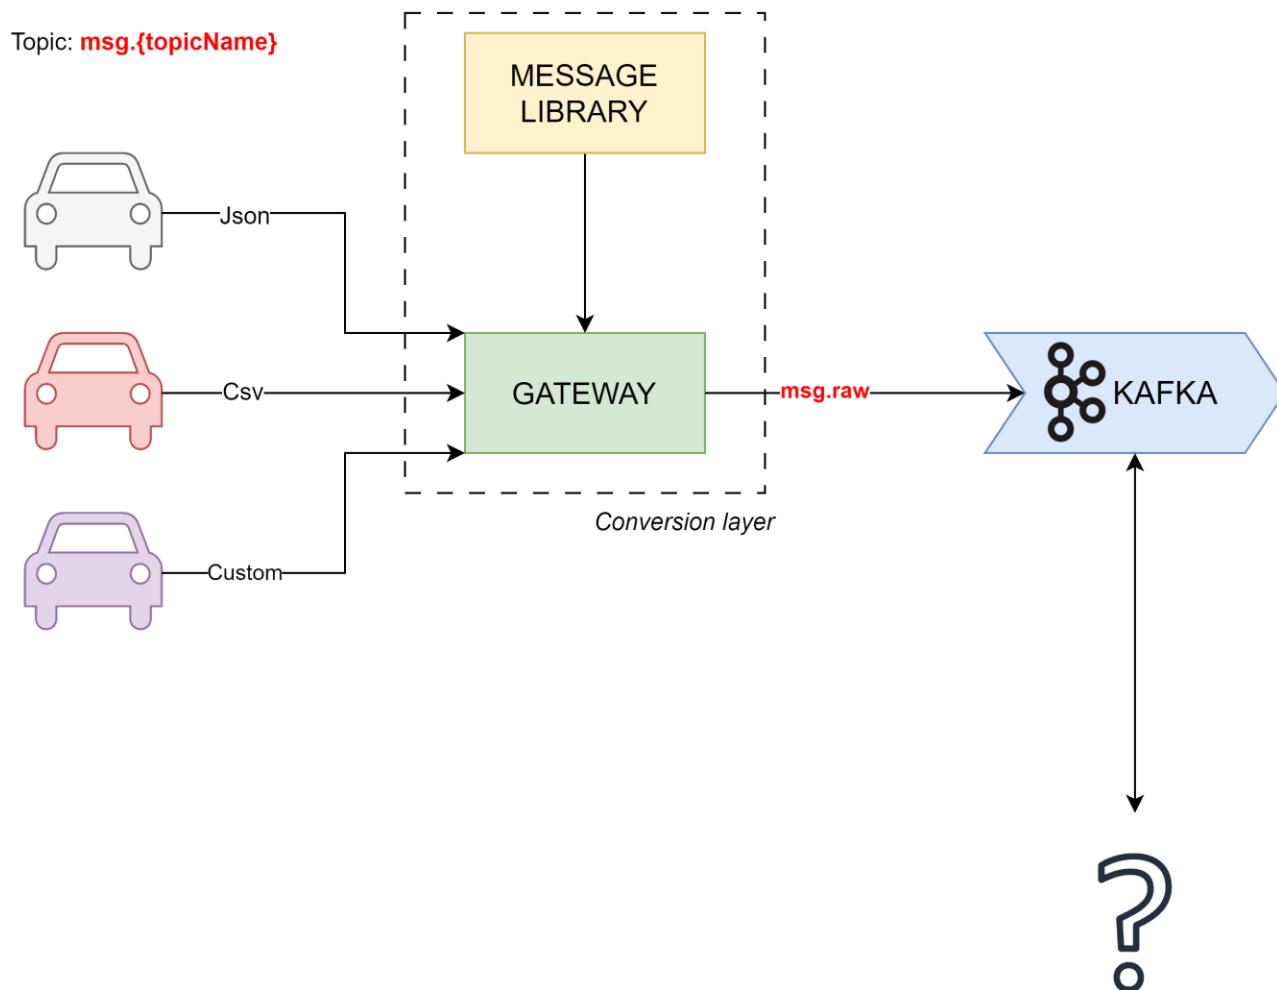
## Parallelism

$P = \text{number of partitions}$ ,  $C = \text{number of consumers}$

$P < C$



# Architecture so far



**What have we achieved?**  
Tracking and logs cars, just a stateless elaboration

**Not enough!**  
We need to remember devices state to get *geofencing*

# Stateless VS Stateful

---

STATELESS	STATEFUL	HYBRID
Process each event independently	Maintains historical data	Combines both trade-offs
No memory of prior events	Complex computations, pattern recognition	Balance low latency with analytics
Suitable for low-latency tasks	Fraud detection, recommendation systems	Choose based on specific use case needs

# Apache Flink

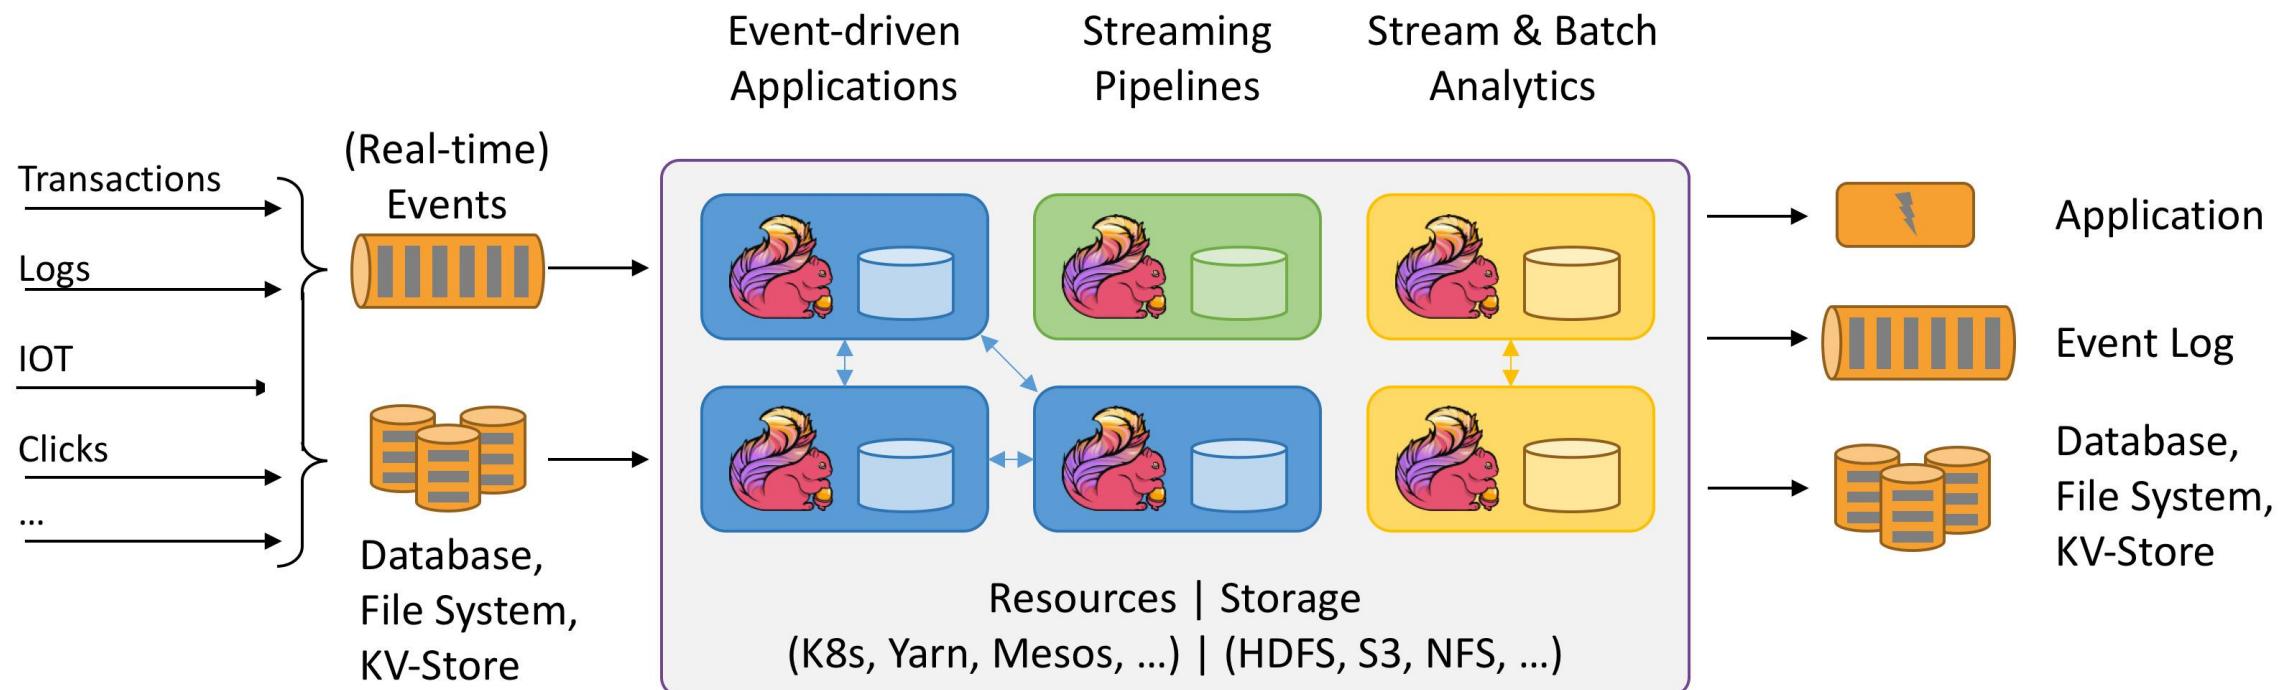
## Processed



Framework for **stateful** computations over *unbounded* and bounded data streams

# Apache Flink

## Framework...



# Apache Flink

---

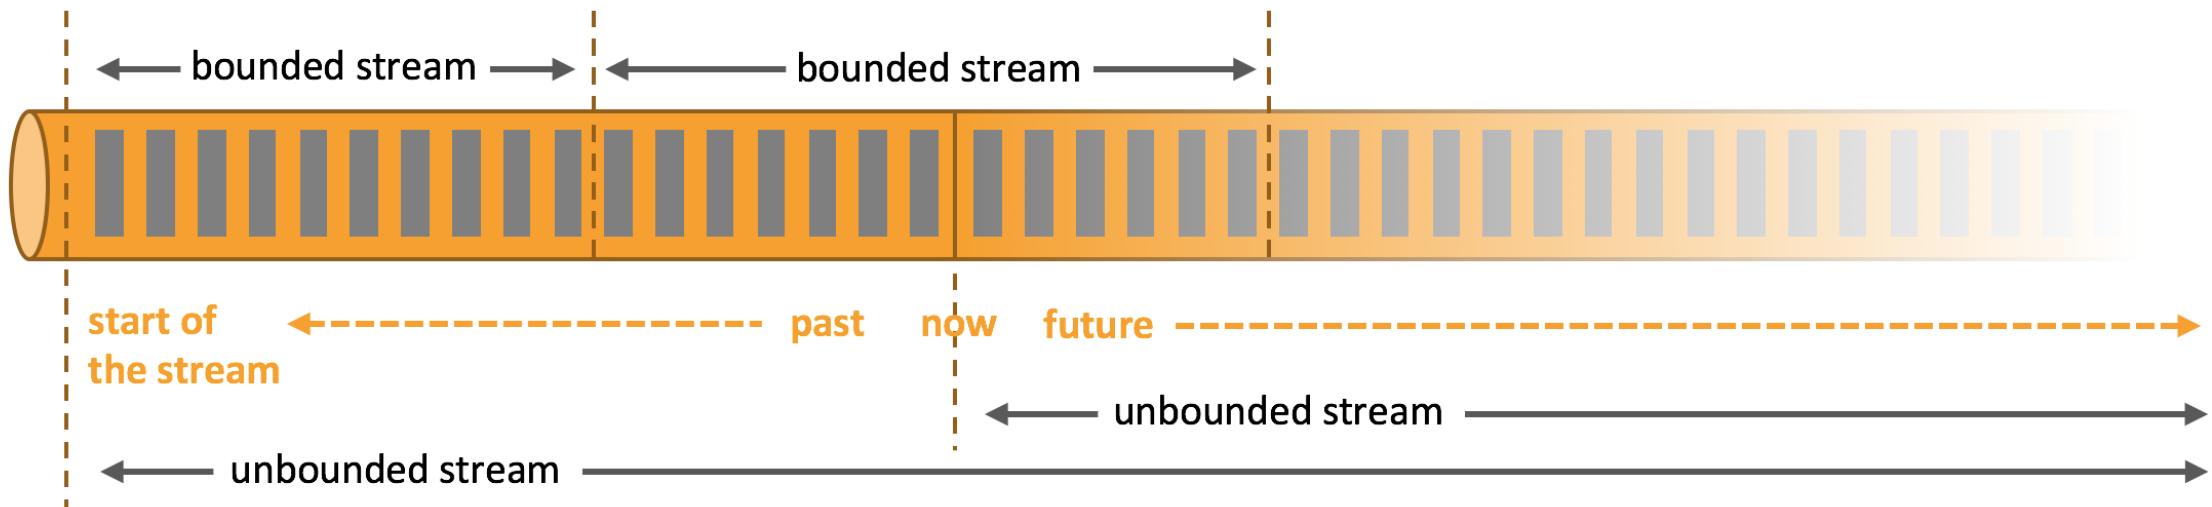
**... for stateful computations...**

Examples of stateful operations:

- Application searching for certain event patterns → sequence of events encountered so far
- Aggregating events per hour → pending aggregates
- Training a machine learning model → current version of the model parameters

# Apache Flink

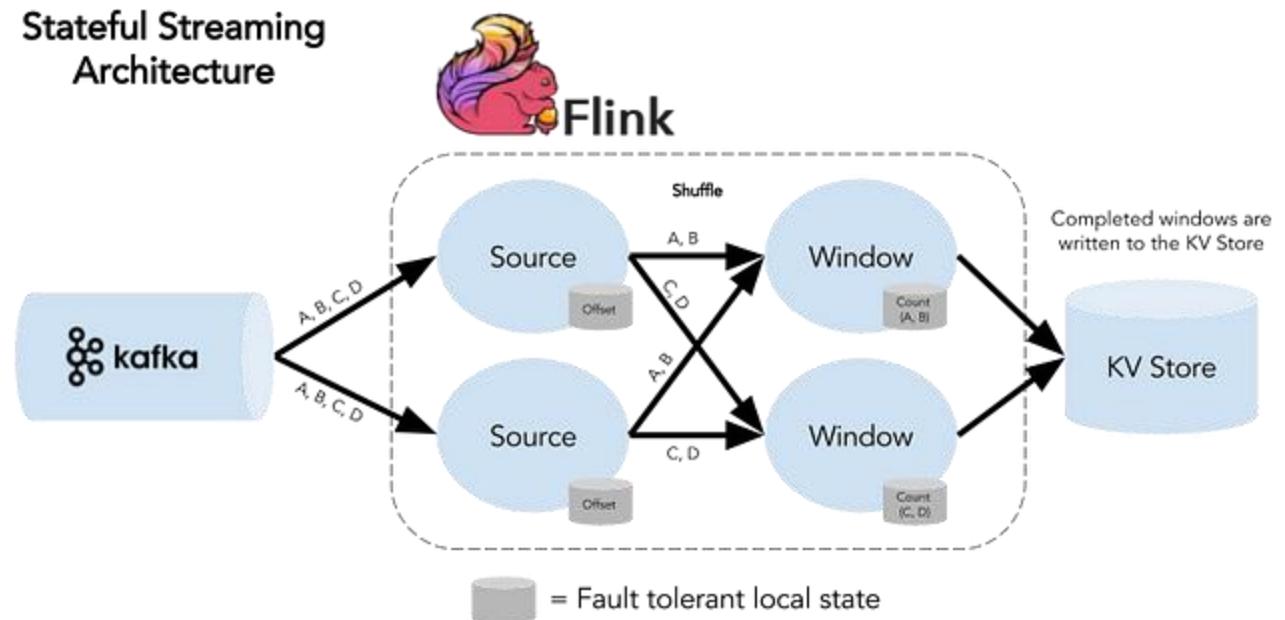
... over *unbounded* and *bounded* data streams



# Apache Flink

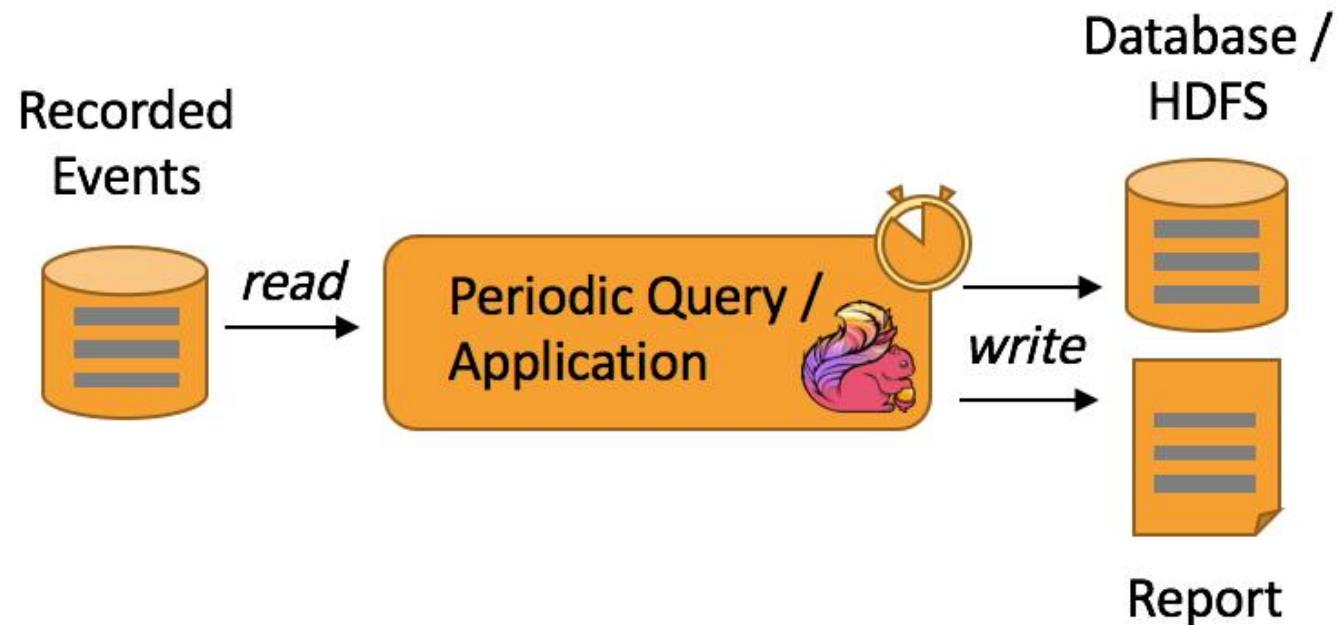
## Features & advantages

- Low latency
- Exactly-once processing
- Stateful processing
- Event time processing
- Fault tolerance
- Scalability
- Integration with Kafka
- Stream-Batch unification



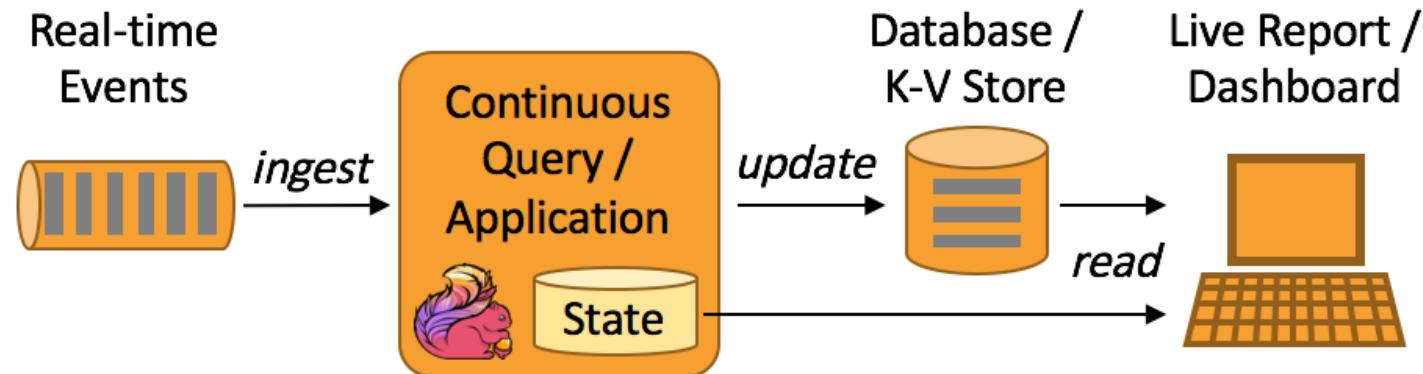
# Apache Flink

## Batch processing



# Apache Flink

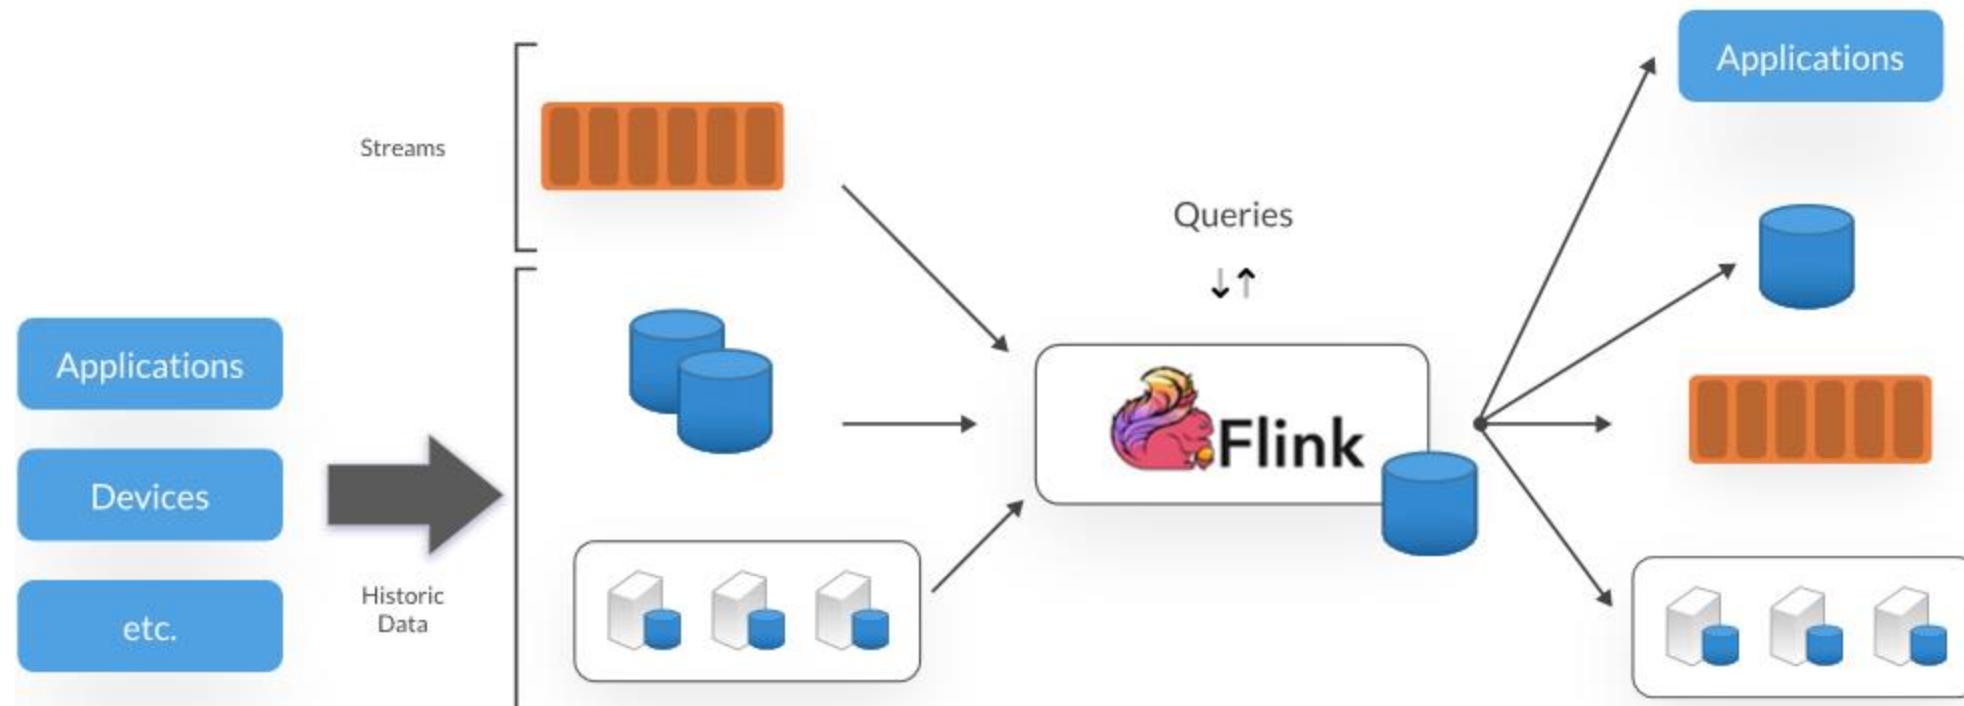
## Stream processing



# Apache Flink

## Processed

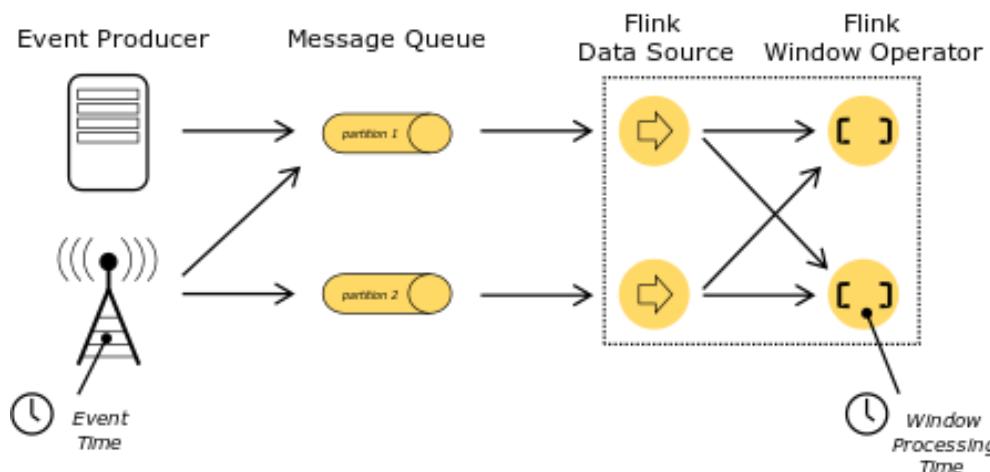
State on data-streams → Stateful → Remember device's state



# Apache Flink

## Timely stream processing

Extension of stream processing in which **time** plays some role in the computation



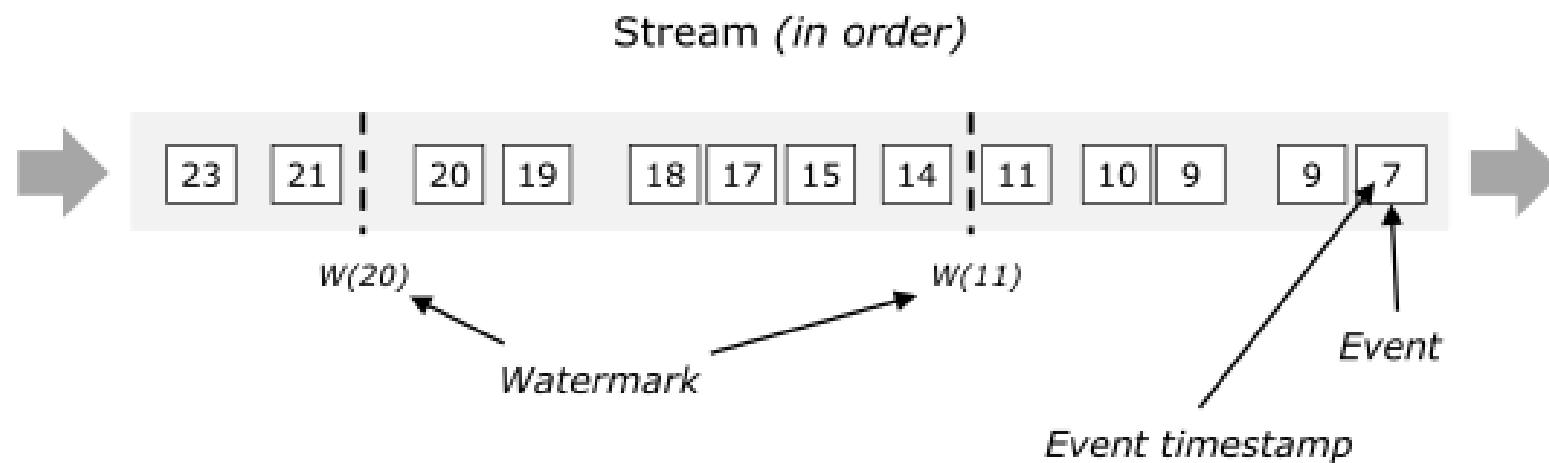
**Processing time:** refers to the system time of the machine that is executing the operation

**Event time:** the time when each individual event occurred

# Apache Flink

## Watermark and Time windows

A **Watermark** declares that event time has reached time  $t$  in that stream, meaning that there should be no more elements from the stream with an event time  $t' \leq t$



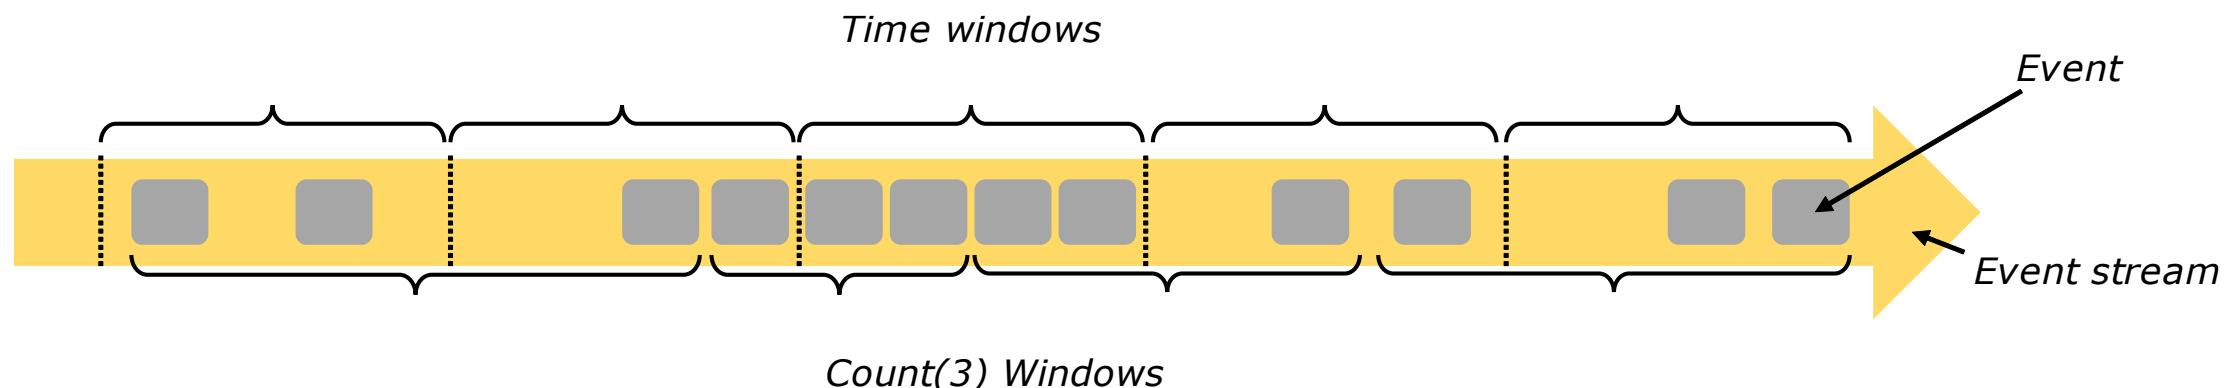
E.G.  $W(11) \rightarrow$  I'll process only records with event time  $\geq 11$

# Apache Flink

## Watermark and Time windows

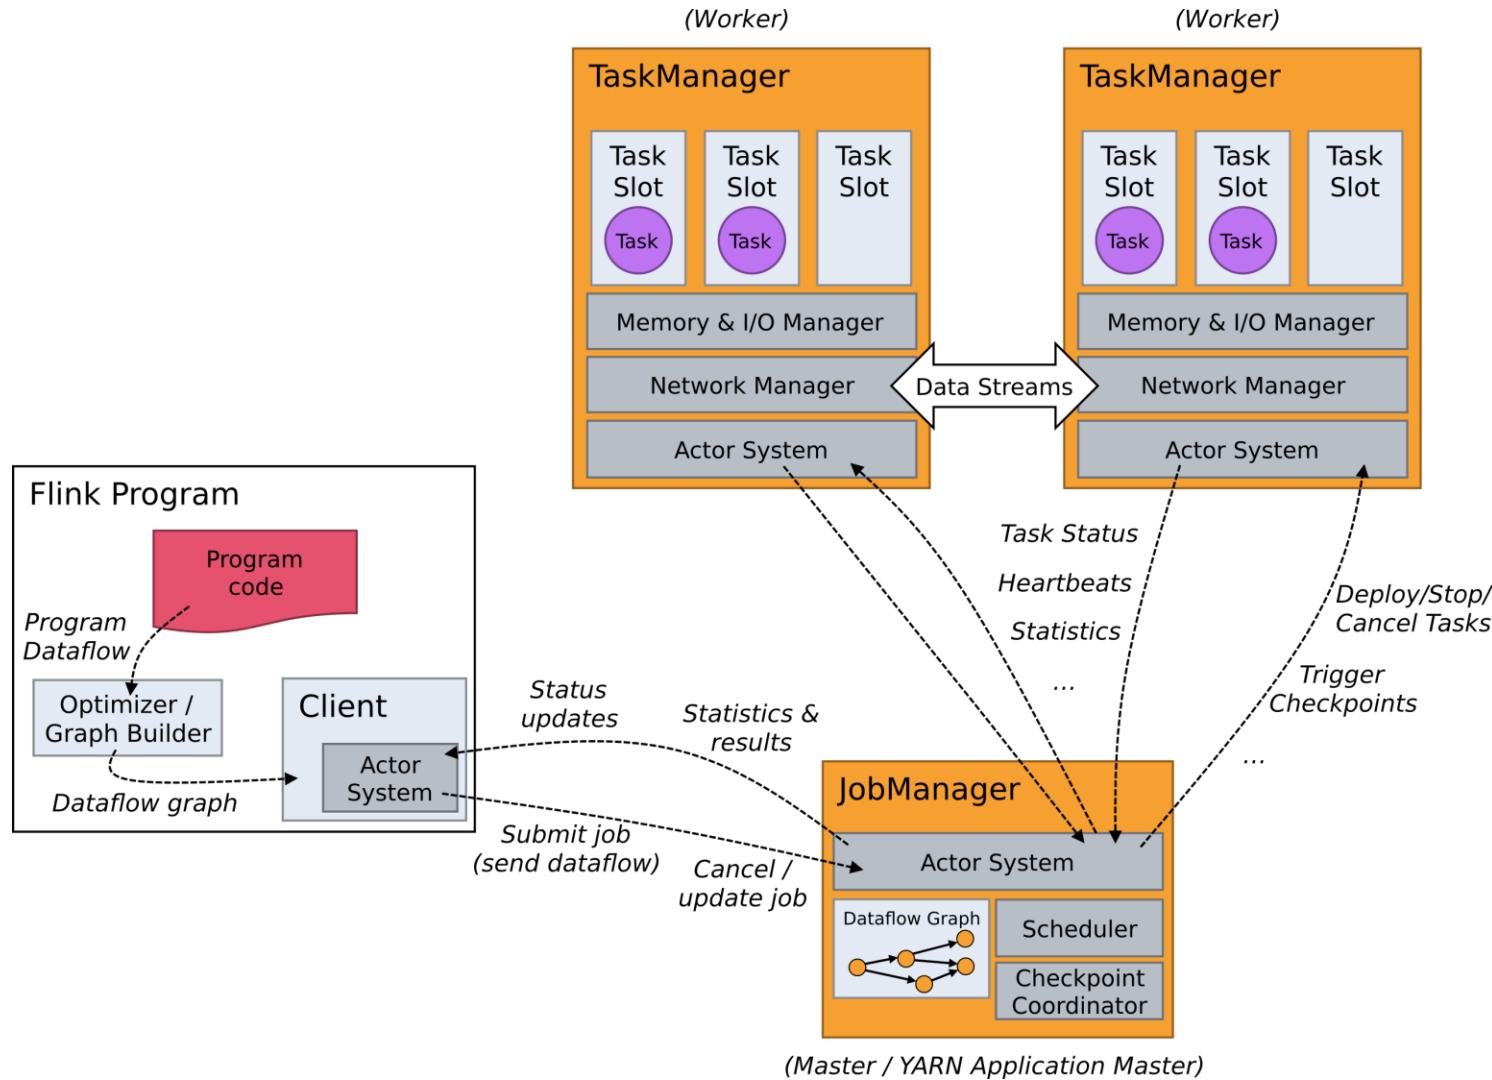
A **Window** can be:

- Time driven or data driven
- Tumbling, sliding or session



# Apache Flink

## Components: JobManager and TaskManager

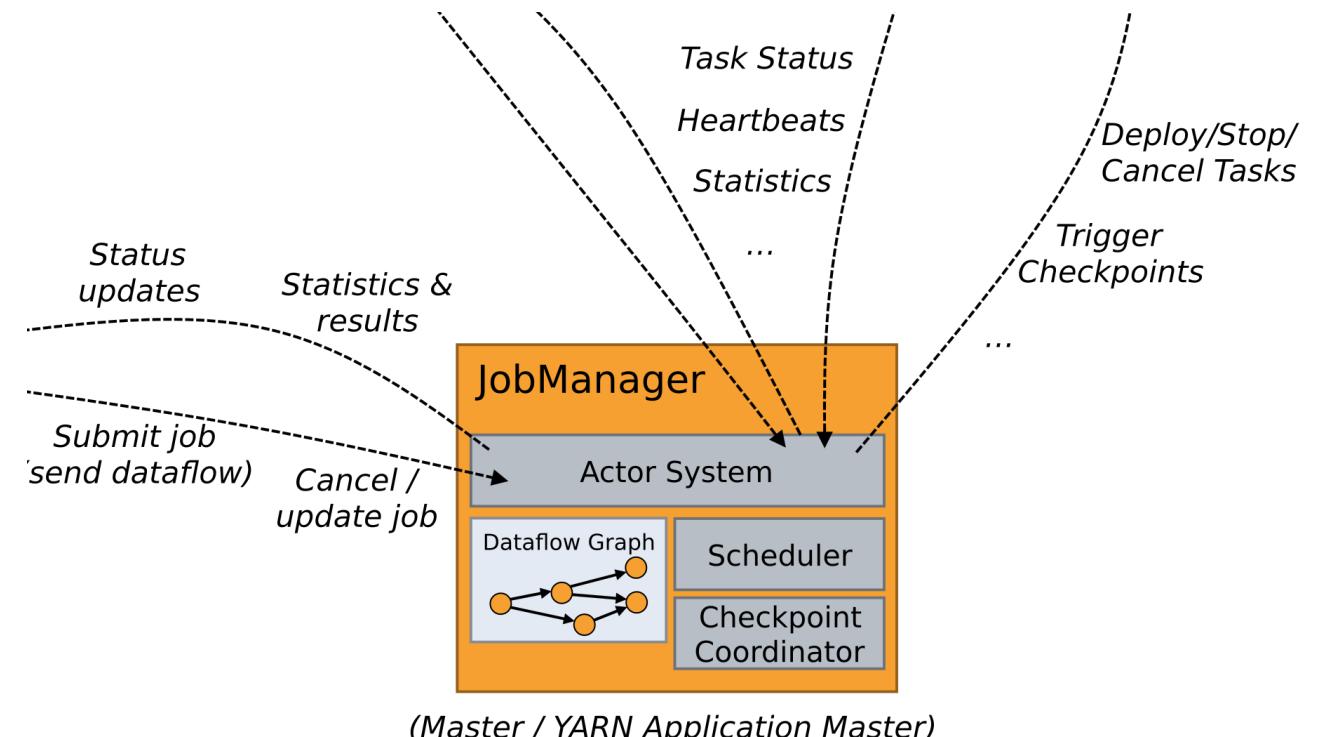


# Apache Flink

## Components: JobManager and TaskManager

### JobManager

In charge of coordinating and distributing execution of Flink application.

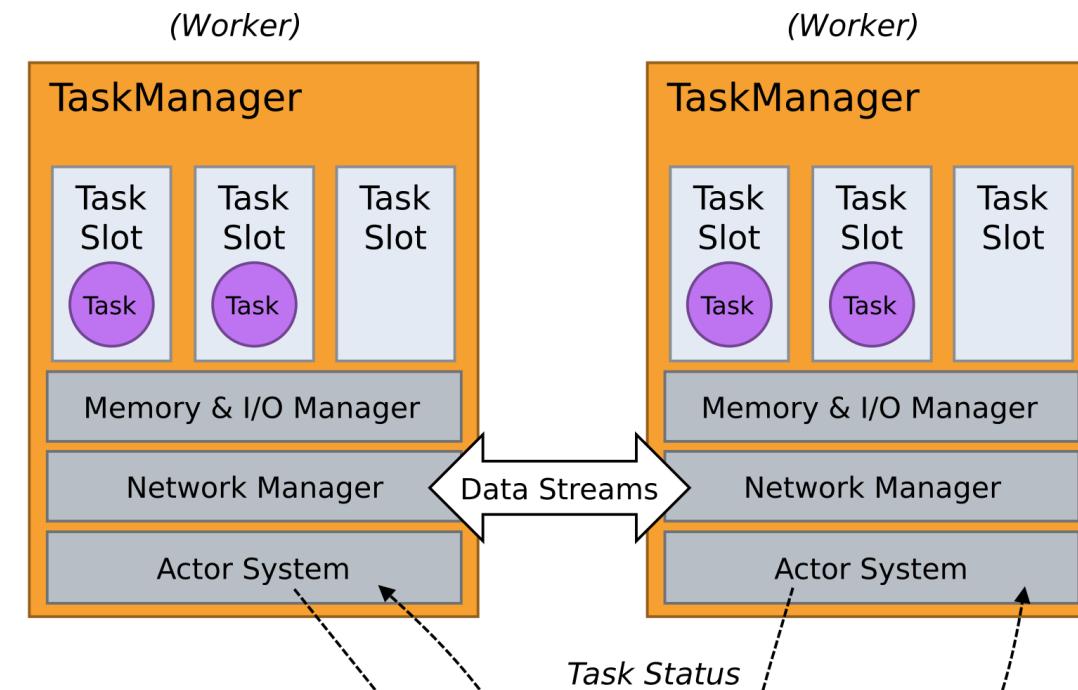


# Apache Flink

## Components: JobManager and TaskManager

### TaskManagers

Execute tasks of a dataflow, and buffer and exchange the data streams.

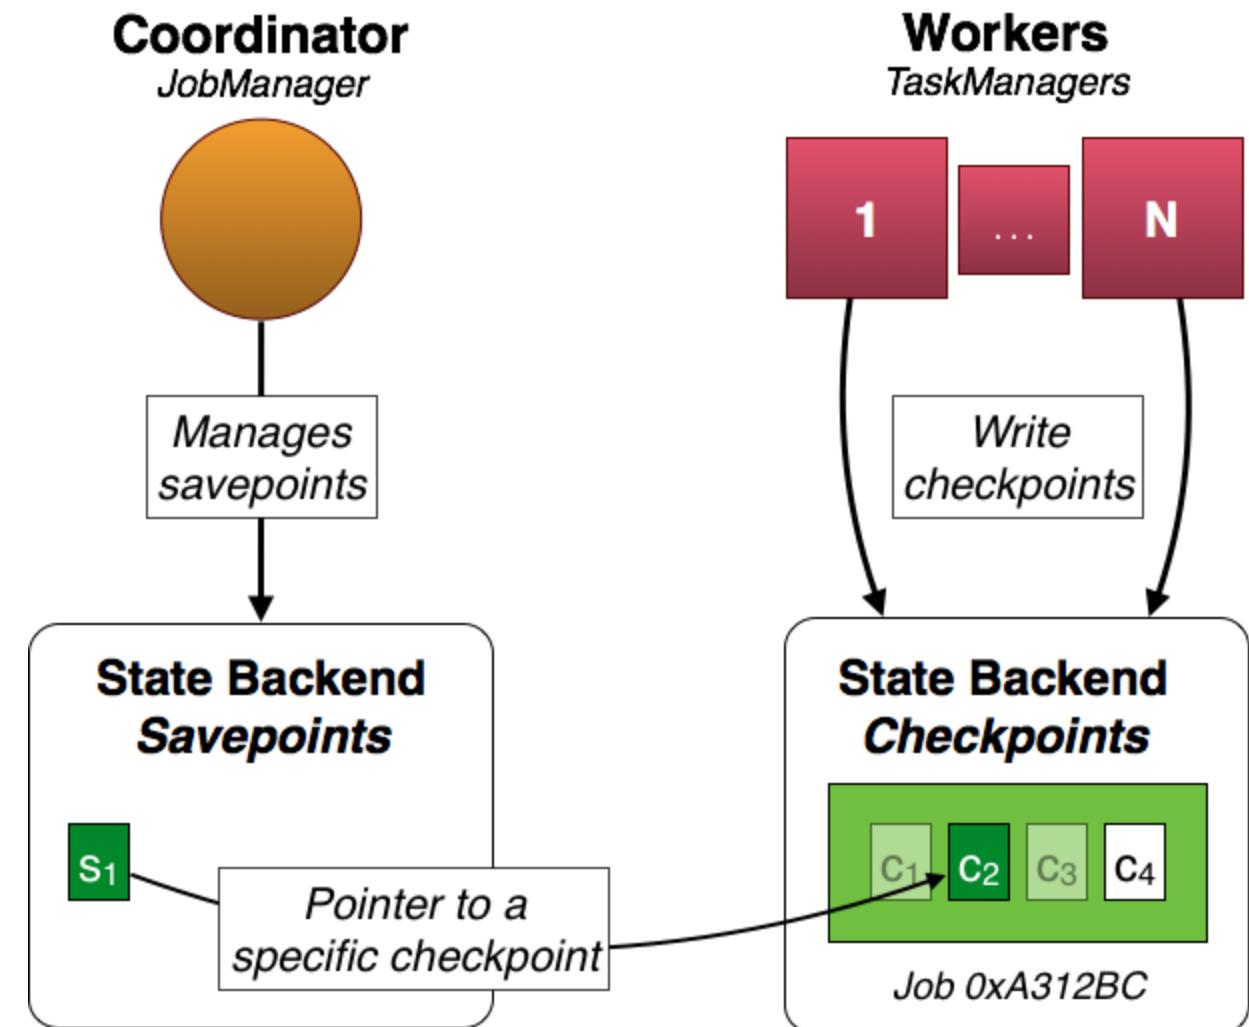


# Apache Flink

## Savepoint - Checkpoint

**Checkpoint:** managed by Flink, relied upon for failure recovery, triggered often.  
Lightweight and fast.

**Savepoint:** planned manual operations, as for example an update of Flink version.



# Apache Flink

---

## Checkpoint recovery

Checkpointing needs:

- Persistent data source that can replay records for a certain amount of time  
*Kafka, RabbitMQ or file systems (HDFS, S3, Ceph)*
- Persistent storage for state  
*HDFS, S3, Ceph..*
- Explicit activation

# Apache Flink

## SQL Abstraction

SQL

High-level Language

Table API

Declarative DSL

DataStream / DataSet API

Core APIs

Stateful Stream Processing

Low-level building block  
(streams, state, [event] time)

# Apache Flink

## SQL Example

```
{  
    "car":{  
        "vin":"PMWR7SPNJAMP38047",  
        "manufacturer":"Acura",  
        "model":"500X",  
        "fuelType":"Compressed Natural Gas",  
        "color":"Beige"  
    },  
    "lat":45.51322,  
    "lon":9.30374,  
    "type":"MOVE"  
}
```

```
{  
    "car":{  
        "vin":"PELP3NP5CPPZ00076",  
        "manufacturer":"Nissan",  
        "model":"G9",  
        "fuelType":"Compressed Natural Gas",  
        "color":"Red"  
    },  
    "lat":45.51323,  
    "lon":9.30381,  
    "type":"MOVE"  
}
```

```
public static void main(String[] args) throws Exception {  
    ...  
  
    String createSourceDDL = "CREATE TABLE JsonSource (\n" +  
        "    car ROW<vin STRING, manufacturer STRING, model STRING, fuelType STRING, color STRING>, \n" +  
        "    lat DOUBLE, \n" +  
        "    lon DOUBLE, \n" +  
        "    type STRING\n" +  
        ") WITH (\n" +  
        "    'connector' = 'kafka', \n" +  
        "    'topic' = 'your_topic', \n" +  
        "    'properties.bootstrap.servers' = 'localhost:9092', \n" +  
        "    'properties.group.id' = 'group_id', \n" +  
        "    'format' = 'json', \n" +  
        "    'json.fail-on-missing-field' = 'false'\n" +  
        ")";  
  
    String query = "SELECT car.vin, car.manufacturer, car.model, car.fuelType, car.color, lat, lon, type FROM  
JsonSource";  
  
    tableEnv.executeSql(createSourceDDL);  
    Table result = tableEnv.sqlQuery(query);  
    tableEnv.toAppendStream(result, Row.class).print();  
}
```

# Apache Flink

## Integration with Apache Kafka

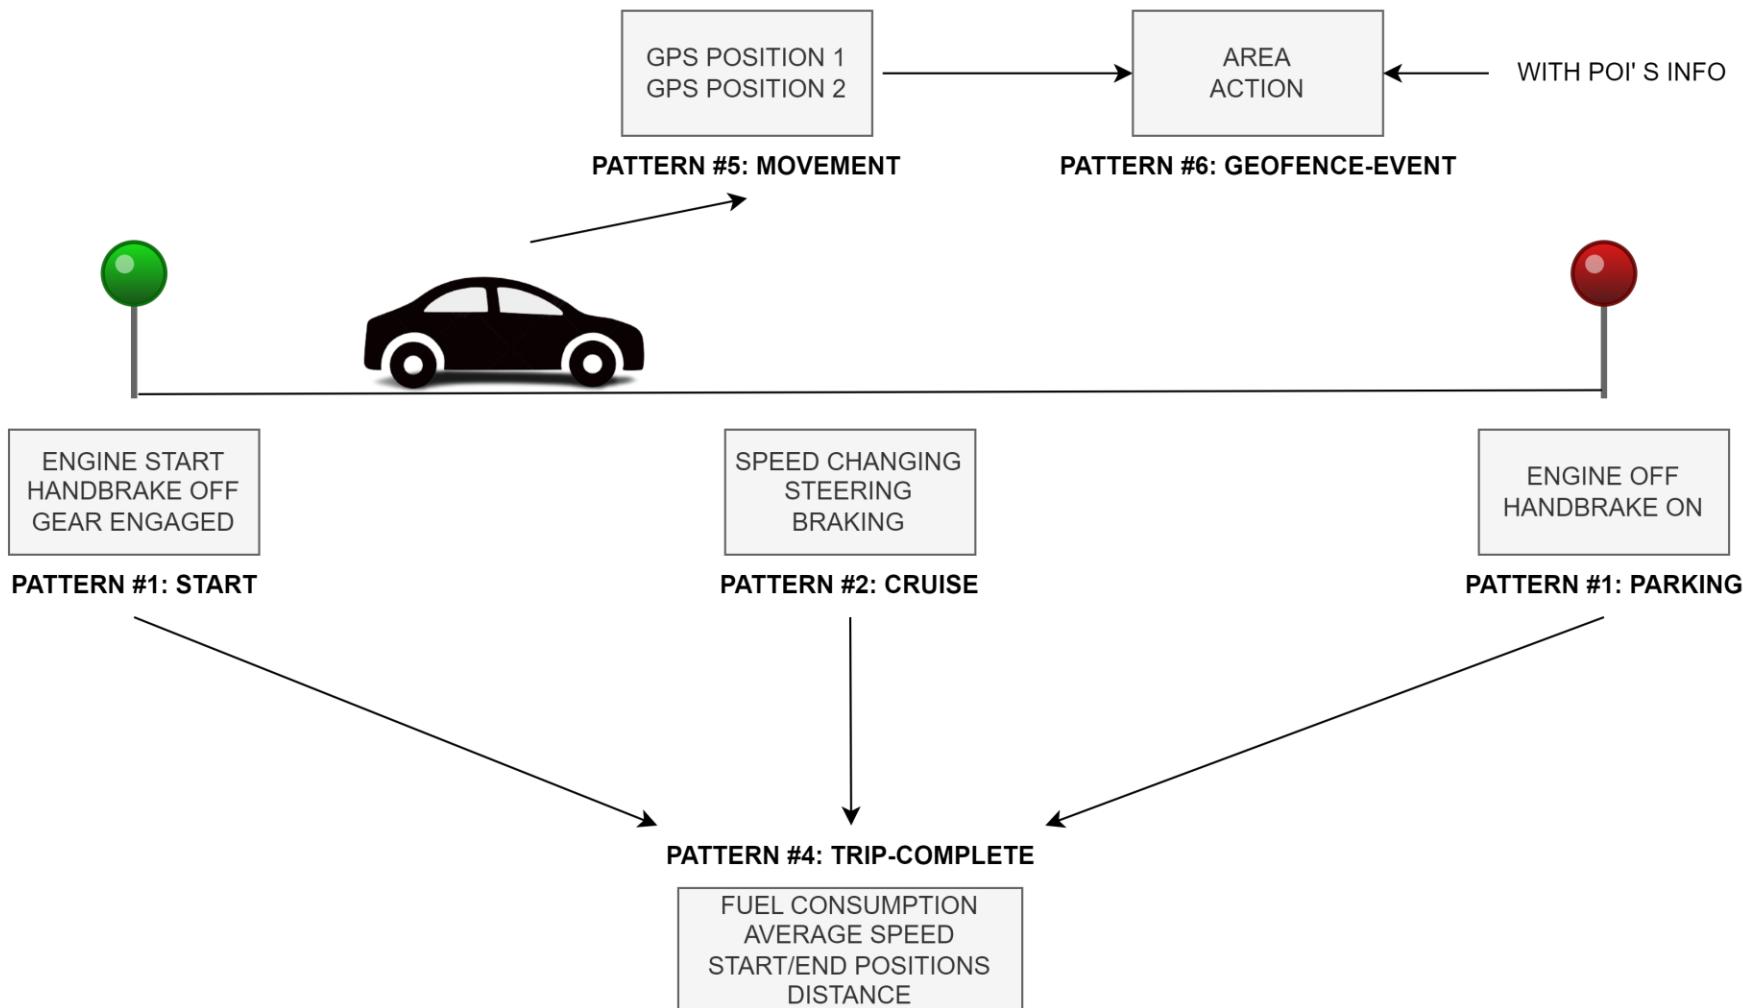
```
KafkaSource<String> source = KafkaSource.<String>builder()
    .setBootstrapServers(brokers)
    .setTopics("input-topic")
    .setGroupId("my-group")
    .setStartingOffsets(OffsetsInitializer.earliest())
    .setValueOnlyDeserializer(new SimpleStringSchema())
    .build();

env.fromSource(source, WatermarkStrategy.noWatermarks(), "Kafka Source");
```

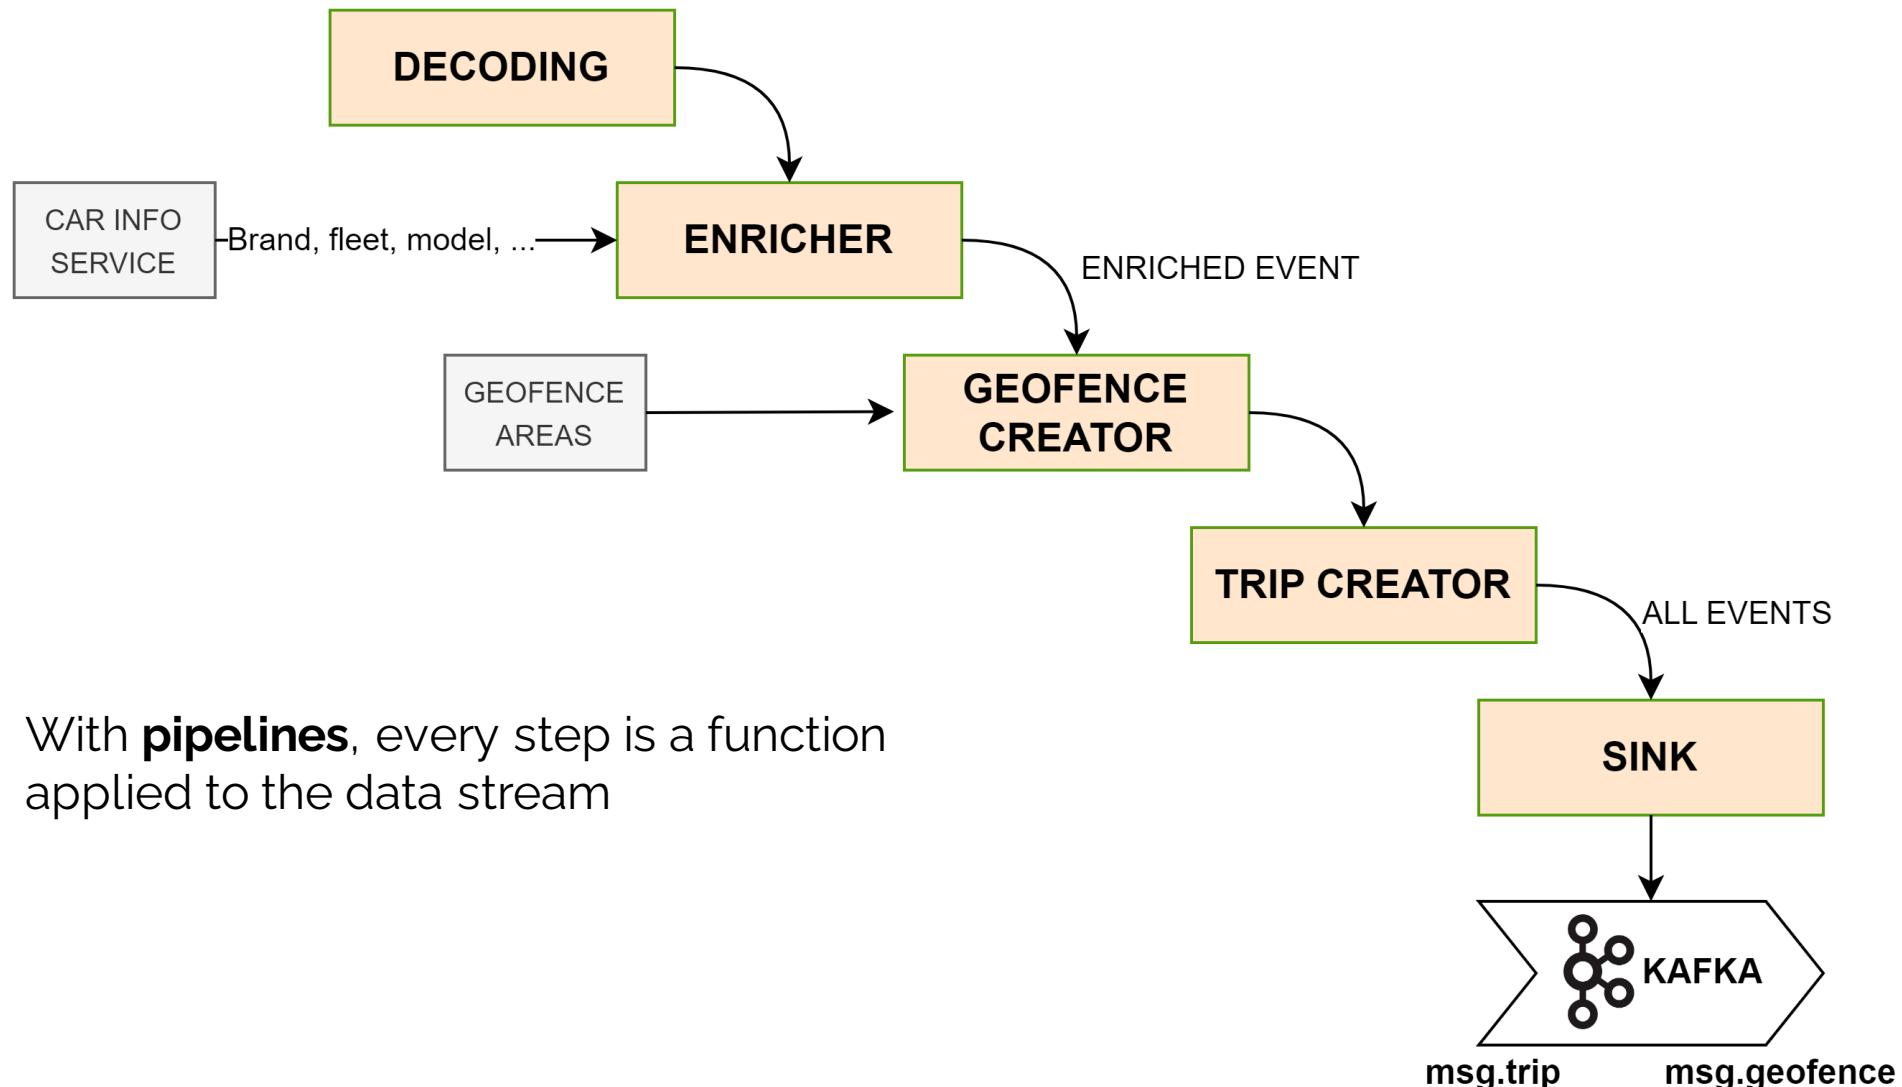
# Use case

## Meaning pattern

A set of data that, once aggregated, forms more structured information

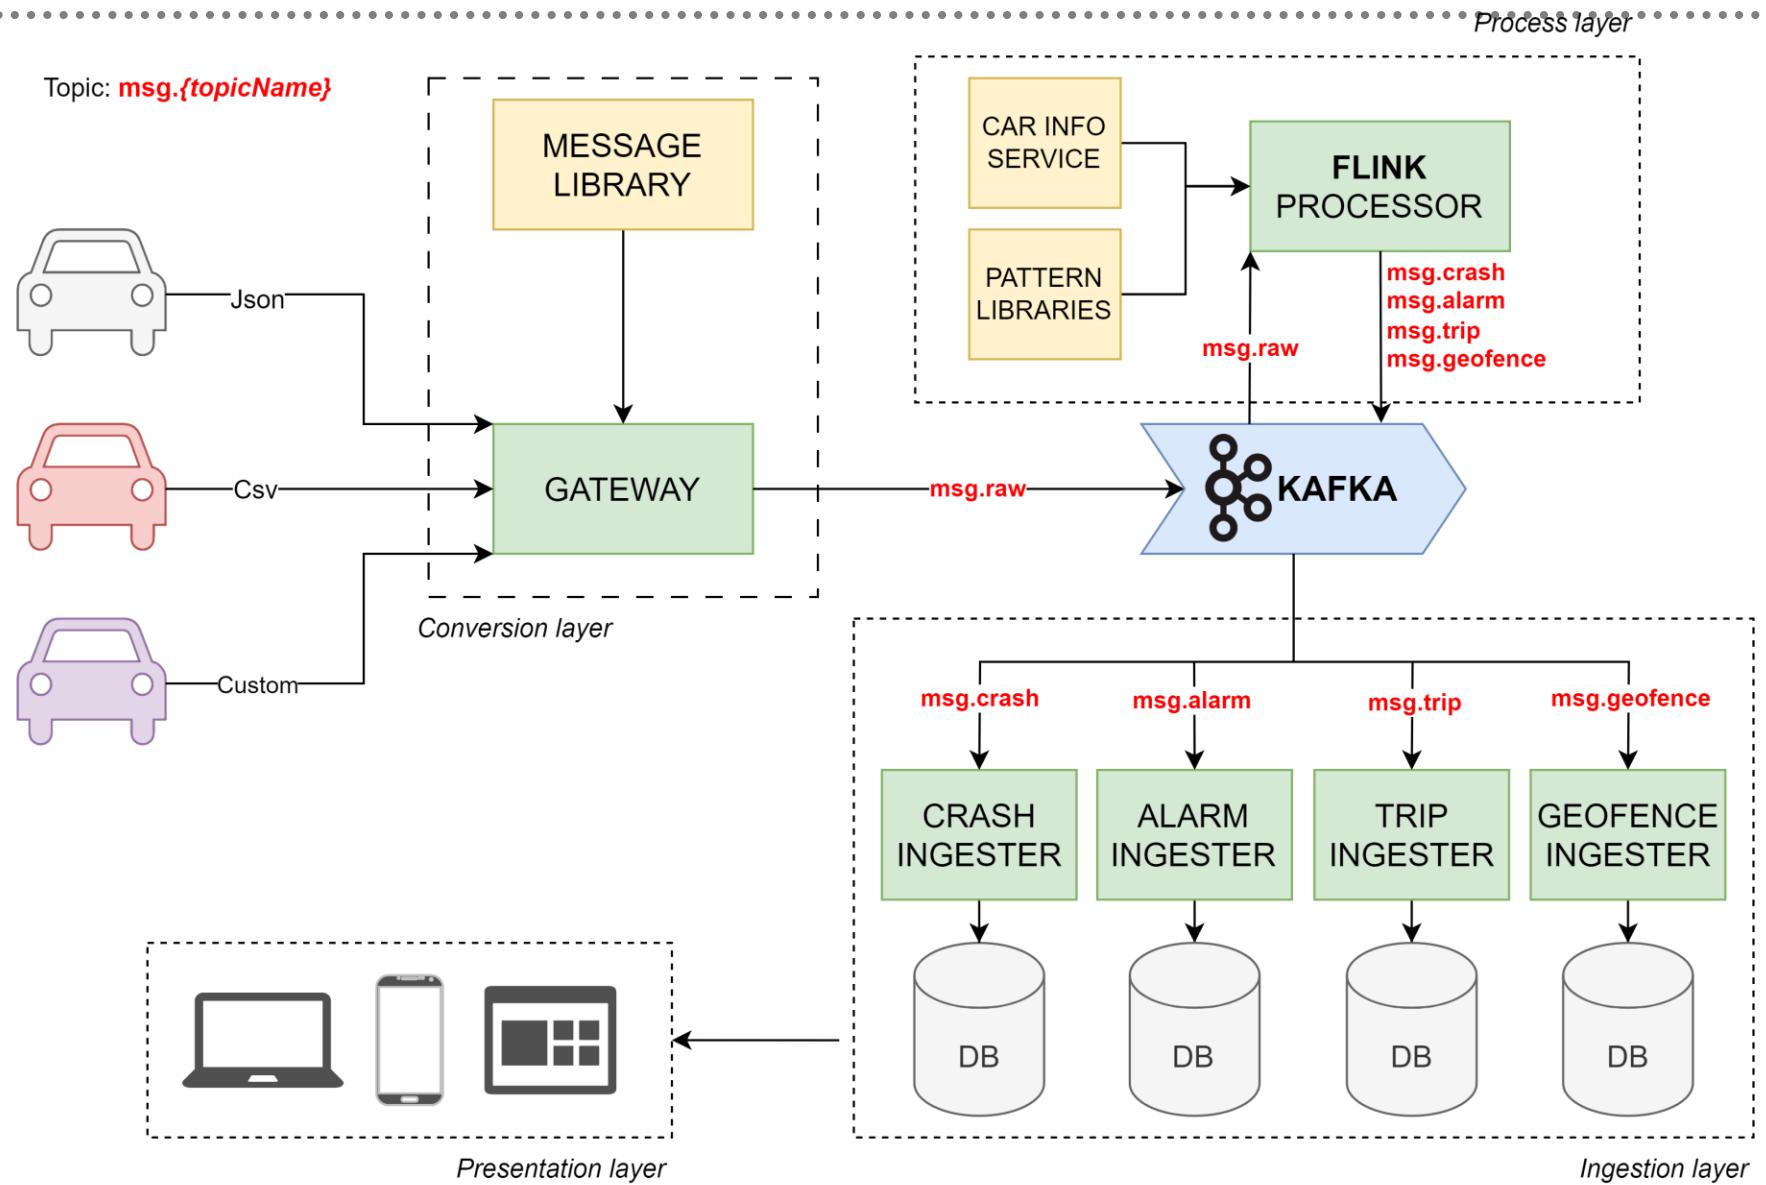


# Build of a pattern

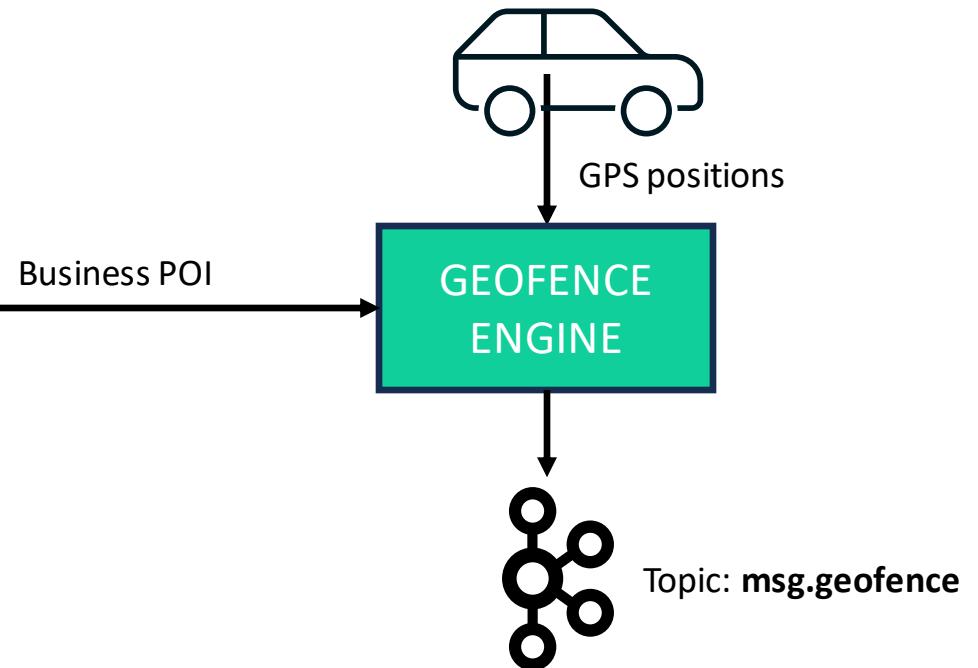
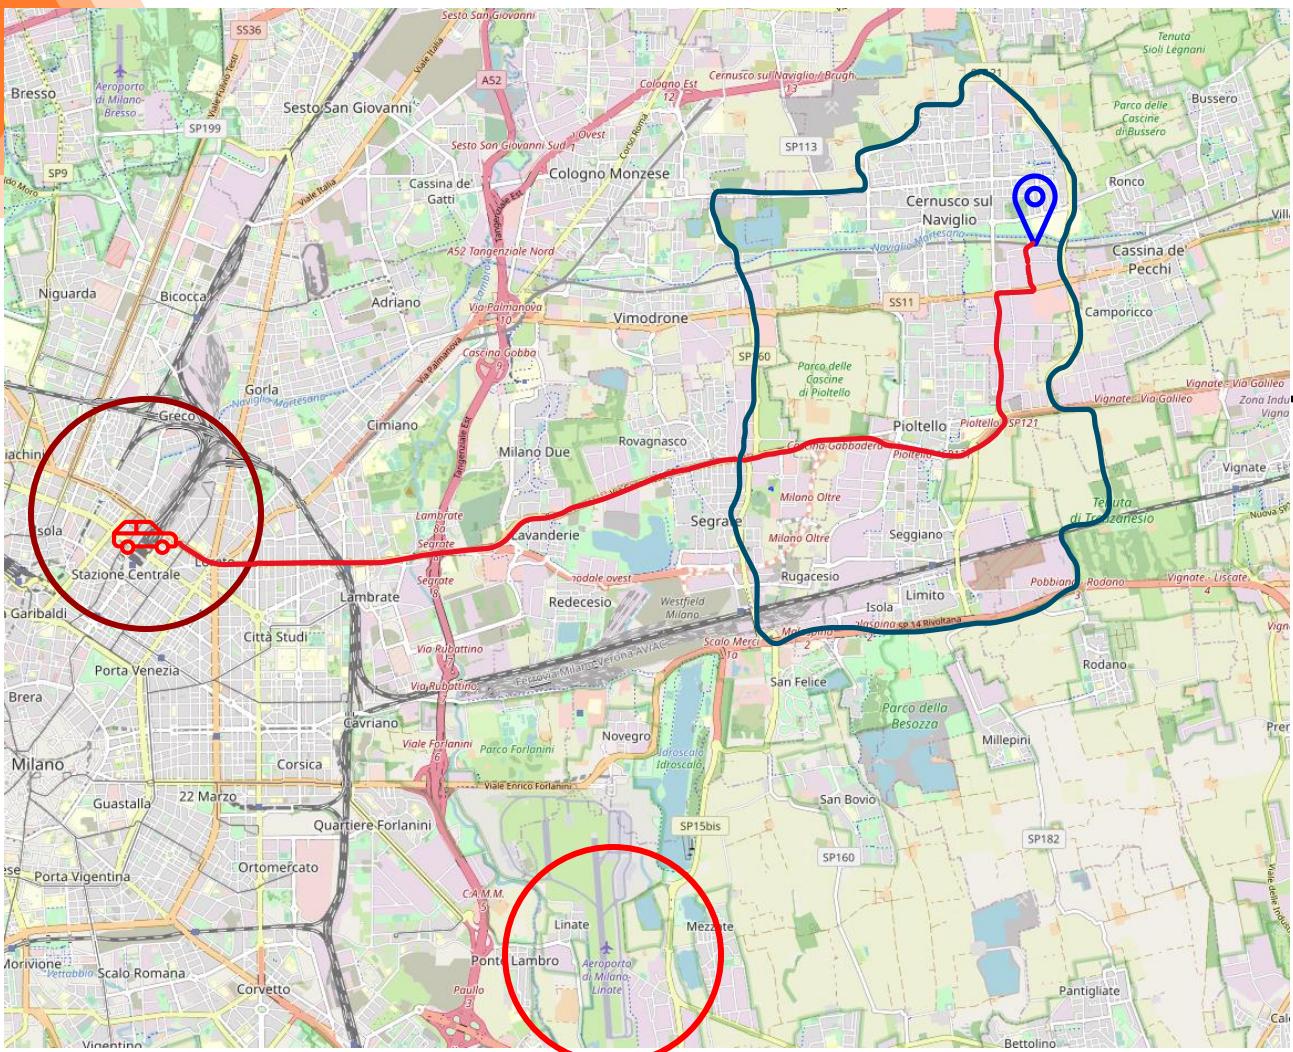


# Architecture so far #2

We refer to whole architecture as  
“Geofence engine”



# “Value”: from raw data to geofencing



{device: 1 , action: start, poi: milano\_centrale}

{device: 1 , action: exit, poi: milano\_centrale}

...

{device: 1 , action: enter, poi: office\_area}

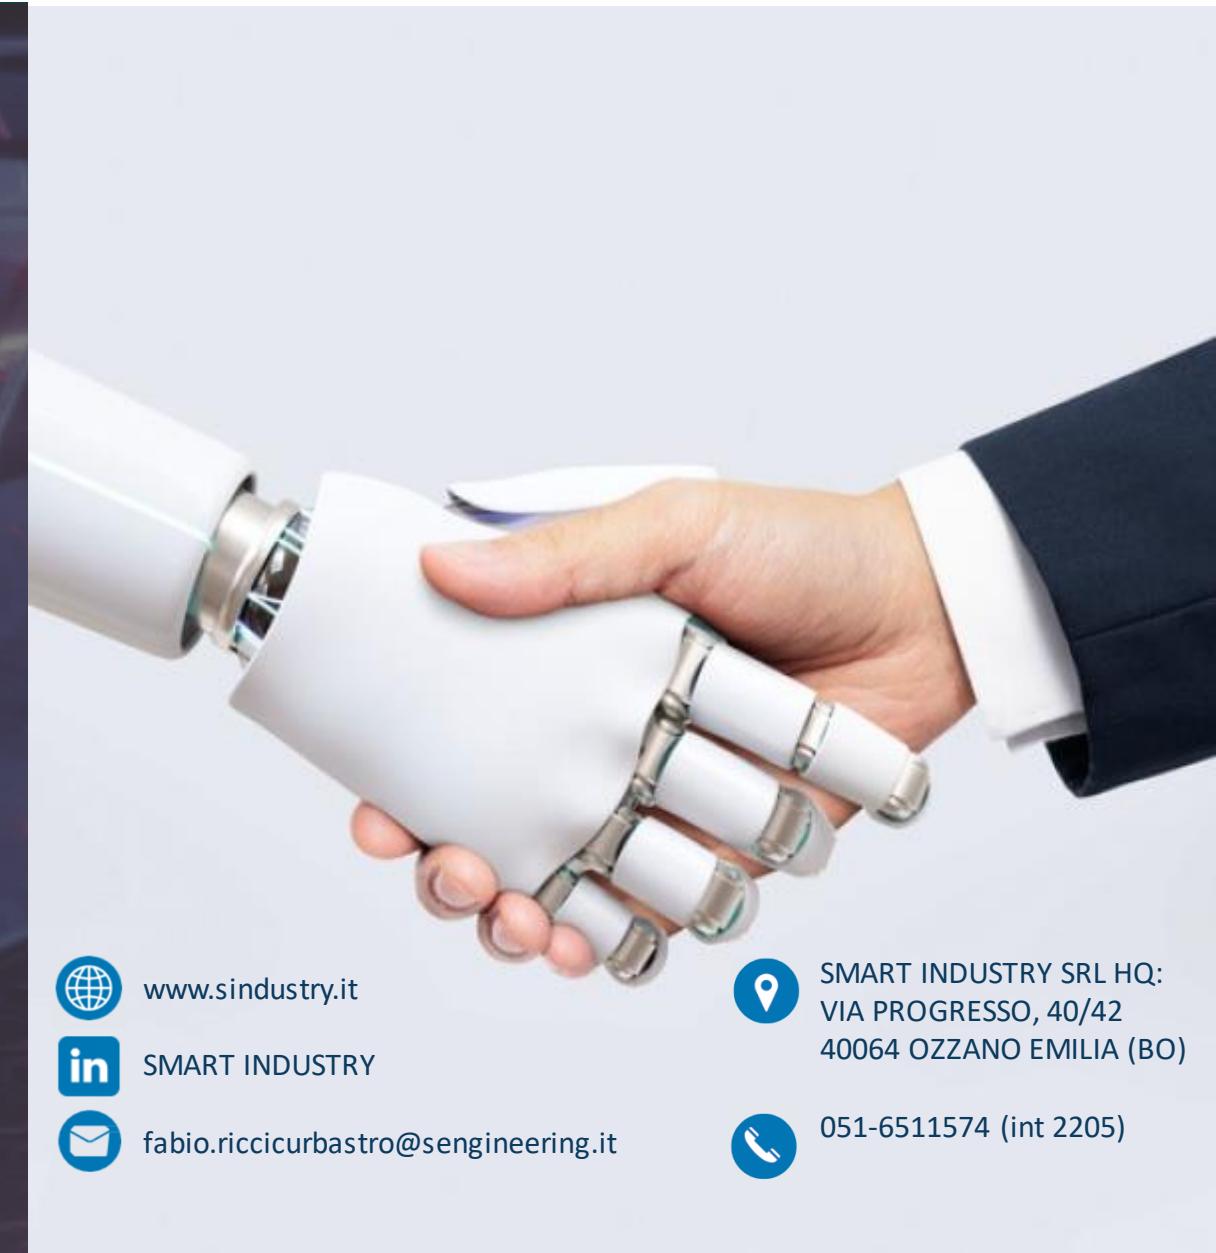
{device: 1 , action: end, poi: headquarter}

Questions about collaborations or opportunities?

Feel free to ask!



# REAL TIME VEHICLE MONITORING AND ANALYSIS



[www.sindustry.it](http://www.sindustry.it)



SMART INDUSTRY



[fabio.riccicurbastro@sengineering.it](mailto:fabio.riccicurbastro@sengineering.it)



SMART INDUSTRY SRL HQ:  
VIA PROGRESSO, 40/42  
40064 OZZANO EMILIA (BO)



051-6511574 (int 2205)